

Technical report for “Conformal prediction for exponential families and generalized linear models”

Daniel J. Eck

February 7, 2019

Abstract

In this supplementary materials document we provide all of the code to reproduce the analyses which appear in the paper Eck et al. [2019] and the R package Eck [2018]. Several additional analyses are provided which illustrate the advantages of conformal prediction regions. In particular, we provide evidence that the parametric conformal prediction region, developed in Eck et al. [2019], compares favorably to other prediction regions even when the parametric model is misspecified.

Contents

1	Introduction/Summary of simulation results	2
2	Illustrative example from the README file	3
3	Extension of the illustrative example from the README file	8
3.1	Diagnostic measures	8
4	Additional Gamma simulations	16
4.1	Simulations	16
4.2	Results	41
5	Gamma-Gaussian model misspecification	48
5.1	Simulations	48
5.2	Results	68
6	Linear Regression Simulations	75
6.1	Simulations	75
6.2	Results	80
7	Example plots of prediction regions	82
7.1	Plots corresponding to Section 4	82
7.2	Plots corresponding to Section 5	91
7.3	Plots corresponding to Section 6	99
8	Depiction of conformal prediction regions that appears in Eck et al. [2019]	103

The following R packages are required in order to replicate the calculations within this document.

```
library(parallel)
library(MASS)
library(statmod)
library(HDInterval)
library(conformal.glm) ## https://github.com/DEck13/conformal.glm
library(conformalInference) ## https://github.com/ryantibs/conformal
```

We set the error tolerance to be $\alpha = 0.10$ for all prediction regions unless otherwise noted.

1 Introduction/Summary of simulation results

This manuscript (and corresponding .Rnw file) provides all of the code to reproduce the analyses which appear in the paper Eck et al. [2019], the README file in the R package Eck [2018], and this document. Several additional analyses to those presented in Eck et al. [2019] are provided in this document. Of particular interest, we investigate the performance of the parametric conformal prediction regions under model misspecification. Specifically, we focus on settings where the underlying data is generated via a Gamma distribution and parametric prediction regions are obtained using a cubic regression model assuming homoscedastic normal errors. The cubic fit is chosen because it is intuitive and it fits the Gamma data better than a simple linear regression model or a quadratic model.

Our goal in this manuscript is to demonstrate the advantages and disadvantages of our parametric conformal prediction region [Eck et al., 2019] compared with the nonparametric conformal prediction region [Lei and Wasserman, 2014], the least squares (LS) conformal prediction region [Lei et al., 2018] obtained from conformalized residual scores, the least squares locally weighted (LSLW) conformal prediction region [Lei et al., 2018, Section 5.2] obtained from conformalized locally weighted residual scores, and the highest density (HD) region. In analyses with model misspecification, the parametric, LS, and LSLW conformal prediction regions and HD prediction region are constructed under the misspecified model. The binning used to construct the parametric and nonparametric conformal prediction regions follows the bin width asymptotics of Lei and Wasserman [2014].

We find that the parametric conformal prediction region performs well even when the model is misspecified. By construction, this region, along with the nonparametric conformal prediction region, maintains finite-sample local validity with respect to binning. These conformal prediction regions therefore achieve finite-sample marginal validity. The guarantee of finite-sample marginal and local validity are noted benefits of these conformal prediction regions [Lei and Wasserman, 2014, Eck et al., 2019]. However, the parametric and nonparametric conformal prediction regions are visually very different, as seen in Section 7, and give different prediction errors at small to moderate sample sizes. We see that the parametric conformal prediction region adapts naturally to the data when the model is correctly specified or modest deviations from the specified model are present. On the other hand, the nonparametric conformal prediction region does not adapt well to data obtained from a Gamma regression model or data obtained from a linear regression model with a steep mean function where steepness is relative to the variability about the mean function.

The LS conformal prediction region obtains marginal validity [Lei et al., 2018] but performs poorly when deviations about the estimated mean function are either not symmetric, not constant, or both. When heterogeneity is present, the LS conformal prediction region exhibits undercoverage in regions where variability about the mean function is large and overcoverage in regions where variability about the mean function is

small. This conformal prediction region is very sensitive to model misspecification. The LSLW conformal prediction region also obtains marginal validity [Lei et al., 2018, Section 5.2] and it is far less sensitive to model misspecification than the LS conformal prediction region. However, the LSLW conformal prediction region is not appropriate when deviations about an estimated mean function are obviously not symmetric, as evidenced in Section 7.1.

2 Illustrative example from the README file

We provide a gamma regression example with perfect model specification to illustrate the performance of conformal predictions when the model is known and the model does not have additive symmetric errors. We also compare conformal prediction regions to the oracle highest density region under the correct model. This example is included in the corresponding paper Eck et al. [2019] and the README file of the corresponding R package `conformal.glm` [Eck, 2018].

```
alpha <- 0.10
n <- 500
shape <- 2
beta <- c(1/4, 2)

set.seed(13)
x <- matrix(runif(n), ncol = 1)
rate <- cbind(1, x) %*% beta * shape
y <- rgamma(n = n, shape = shape, rate = rate)
data.readme <- data.frame(y = y, x = x)
colnames(data.readme)[2] <- c("x1")

fit.readme = glm(y ~ x1, family = Gamma, data = data.readme)
summary(fit.readme)

##
## Call:
## glm(formula = y ~ x1, family = Gamma, data = data.readme)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0141  -0.7055  -0.1658   0.3386   1.7546
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.24070     0.02671   9.012  <2e-16 ***
## x1           2.04663     0.10584  19.336  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for Gamma family taken to be 0.4786753)
##
##      Null deviance: 469.94  on 499  degrees of freedom
```

```
## Residual deviance: 257.53 on 498 degrees of freedom
## AIC: 817.06
##
## Number of Fisher Scoring iterations: 6
```

We now compute the parametric conformal prediction region [Eck et al., 2019] and the nonparametric conformal prediction region [Lei and Wasserman, 2014] using the `conformal.glm` function in the `conformal.glm` package. For illustration, we set the number of bins to be equal to 3. Since $X_i \sim U(0, 1)$, we expect for there to be about 167 cases per bin with this choice of 3 bins. This choice for the number of bins is consistent with the bin width asymptotics in Lei and Wasserman [2014]. The number of cores is set to 6, if your computer has fewer than 6 cores at its disposal then you have to reduce the number of cores. This will not effect the final output but it will take longer to run.

```
bins <- 3
system.time(cpred <- conformal.glm(fit.readme,
  nonparametric = TRUE, bins = bins, cores = 6))

##      user  system elapsed
## 140.926    0.215   41.171

paraCI <- cpred$paraconformal
nonparaCI <- cpred$nonparaconformal
```

We now compute the least squares with local weighting (LSLW) conformal prediction region [Lei et al., 2018, Section 5.2] using the `conformal.pred` function in the `conformalInference` package [Tibshirani, 2016]. In this procedure, the residuals from a regression fit are weighted by an estimate of the conditional mean absolute deviation (MAD) of $(Y - \mu(X))|X = x$, as a function of x . For augmented data including a proposed y , we let $\hat{\rho}_y(x)$ denote the estimate of the MAD. The conformal prediction procedure now uses

$$R_{y,i} = \frac{|Y_i - \hat{\mu}_y(X_i)|}{\hat{\rho}_y(X_i)}, \quad (i = 1, \dots, n+1)$$

as a anti-conformity measure. See Section 5.2 in Lei et al. [2018] for more details. The code below implements this method computing residuals from a cubic regression model, the cubic model fits the Gamma data better than a simple linear or quadratic model. We estimate $\hat{\rho}_y$ using a smoothing spline whose parameters are chosen by cross-validation.

```
library(conformalInference)
funs <- lm.funs(intercept = TRUE)
train.fun <- funs$train.fun
predict.fun <- funs$predict.fun

cubic.model <- lm(y ~ x + I(x^2) + I(x^3))
abs.resid <- abs(cubic.model$resid)
smooth.call <- smooth.spline(x, abs.resid,
  nknots = 10)
lambda <- smooth.call$lambda
df <- smooth.call$df
mad.train.fun <- function(x, y, out = NULL) {
```

```

smooth.spline(x[, 1], y, lambda = lambda,
              df = df, nknots = 10)
}
mad.predict.fun <- function(out, newx){
  predict(out, newx[, 1])$y
}
system.time(pl.tibs <- conformal.pred(x = cbind(x, x^2, x^3),
  y = y, x0 = cbind(x, x^2, x^3),
  train.fun = train.fun, predict.fun = predict.fun,
  mad.train.fun = mad.train.fun,
  mad.predict.fun = mad.predict.fun,
  alpha = alpha))

##      user  system elapsed
##  89.826    0.087   89.934

LSLW = cbind(pl.tibs$lo, pl.tibs$up)

```

We now compute the highest density region using the `hdi` function in the `HDInterval` package [Meredith and Kruschke, 2018].

```

library(HDInterval)
betaMLE <- coefficients(fit.readme)
shapeMLE <- as.numeric(gamma.shape(fit.readme)[1])
rateMLE <- cbind(1, x) %**% betaMLE * shapeMLE
HDCI <- do.call(rbind,
  lapply(1:nrow(x), function(j){
    hdi(qgamma, 0.90, shape = shapeMLE, rate = rateMLE[j, 1])
  }))

```

The four prediction regions for this data are depicted below. The top row depicts the parametric conformal prediction region (left panel) and the nonparametric conformal prediction region (right panel). The bin width was specified as $1/3$ for these conformal prediction regions. The bottom row depicts the least squares conformal prediction region (left panel) and the highest density region (right panel). We see that the parametric conformal prediction region is a close discretization of the highest density region, the nonparametric conformal prediction region is quite jagged and unnatural, and the least squares conformal prediction region exhibits undercoverage for small x , exhibits overcoverage for large x , and includes negative response values of large magnitude.

All of the presented prediction regions exhibit close to finite-sample marginal validity and local validity with respect to binning. However, the LSLW conformal prediction region and the HD prediction region do not exhibit finite-sample local validity in the second bin and the HD prediction region does not quite possess finite-sample marginal validity. The parametric conformal prediction region is smallest in size with an estimated area of 2.20. The HD prediction region is a close second with an estimated area of 2.21. LSLW conformal prediction region has an estimated area of 2.57 and The nonparametric conformal prediction region has an estimated area of 2.69. Under correct model specification, the parametric conformal prediction region is similar in performance to that of the highest density prediction region.




Figure 1: Sim setting: $n = 500$, $\text{shape} = 2$, $\text{bins} = 3$. The regions are depicted as follows: parametric conformal prediction region (top-left panel), nonparametric conformal prediction region (top-right panel), LSLW conformal prediction region (bottom-left panel), HD prediction region (bottom-right panel).

```
## parametric conformal prediction region
# estimated area
mean(apply(paraCI, 1, diff))

## [1] 2.198495

# local coverage
```

```

p <- length(beta) - 1
local.coverage(region = paraCI, data = data.readme, d = p,
  bins = bins, at.data = "TRUE")

## [1] 0.9135802 0.9028571 0.9079755

# marginal coverage
local.coverage(region = paraCI, data = data.readme, d = p,
  bins = 1, at.data = "TRUE")

## [1] 0.908

## nonparametric conformal prediction region
# estimated area
area.nonparametric <- function(region){
  if(class(region) != "list"){
    stop("Only appropriate for nonparametric conformal prediction region")
  }
  bins <- length(region); wn <- 1 / bins
  area <- 0
  for(i in 1:bins){
    nonpar.region <- region[[i]]
    area <- area + wn * as.numeric(rep(c(-1,1),
      length(nonpar.region)/2) %*% nonpar.region)
  }
  area
}
area.nonparametric(nonparaCI)

## [1] 2.687667

# local coverage
local.coverage(region = nonparaCI, data = data.readme, d = p,
  nonparametric = "TRUE", bins = bins, at.data = "TRUE")

## [1] 0.9259259 0.9028571 0.9018405

# marginal coverage
local.coverage(region = nonparaCI, data = data.readme, d = p,
  nonparametric = "TRUE", bins = 1, at.data = "TRUE")

## [1] 0.91

## LSLW conformal prediction region
# estimated area
mean(apply(LSLW, 1, diff))

## [1] 2.563743

```

```

# local coverage
local.coverage(region = LSLW, data = data.readme, d = p,
  bins = bins, at.data = "TRUE")

## [1] 0.9197531 0.8914286 0.9509202

# marginal coverage
local.coverage(region = LSLW, data = data.readme, d = p,
  bins = 1, at.data = "TRUE")

## [1] 0.92

## HD region
# estimated area
mean(apply(HDCI, 1, diff))

## [1] 2.207619

# local coverage
local.coverage(region = HDCI, data = data.readme, d = p,
  bins = bins, at.data = "TRUE")

## [1] 0.9012346 0.8800000 0.9079755

# marginal coverage
local.coverage(region = HDCI, data = data.readme, d = p,
  bins = 1, at.data = "TRUE")

## [1] 0.896

```

3 Extension of the illustrative example from the README file

In this extension we investigate the local coverage properties of five prediction regions, the four from the previous section and a least squares (LS) conformal prediction region without local weighting. Local coverage is assessed via a Monte Carlo simulation of size $B = 50$. At each iteration of this simulation a new dataset is generated under the same specifications as those in the README file. We then compute the local coverage probabilities with respect to each bin and also approximate conditional coverage across x . We do not assess coverage properties directly at x for all $x \in (0, 1]$. This corresponds to the notion of conditional validity which cannot be achieved in tandem with an oracle estimation in finite sample settings when distributions are continuous [Lei and Wasserman, 2014, Section 2.2]. What we do instead is we assess local coverage with a binning in x that is much finer than the binning that was used to create the misspecified parametric and nonparametric conformal prediction regions. We form 25 bins of length 0.04.

3.1 Diagnostic measures

In this section we describe the diagnostic measures that we will use to compare conformal prediction regions. The specific diagnostics are prediction error at the observed data, area of the prediction region, finite-sample

marginal coverage, finite-sample local coverage with respect to binning, and finite-sample local conditional coverage.

We will use a modification of average sum of squares error as our measure of prediction error. Our prediction error will average the squared distances of observations outside of the prediction region to the closest boundary of the prediction region. The average will be taken over all observations so that an observation that falls within the prediction region has an error of 0. More formally this prediction error is

$$\text{prediction error} = n^{-1} \sum_{i=1}^n \mathbb{1} \left\{ Y_i \notin C^{(\alpha)}(X_i) \right\} (\min\{|Y_i - a_{i,1}|, |Y_i - b_{i,1}|, \dots, |Y_i - a_{i,m_i}|, |Y_i - b_{i,m_i}|\})^2,$$

where $a_{i,j}$ and $b_{i,j}$ are, respectively, the lower and upper boundaries of possible $j = 1, \dots, m_i$ disjoint intervals forming the prediction region.

The area of each prediction region will be estimated by the average of the upper boundary minus the lower boundary across observed \mathcal{X} , written as

$$\text{area} = n^{-1} \sum_{i=1}^n \sum_{j=1}^{m_i} (b_{i,j} - a_{i,j}).$$

This estimate of area is appropriate in our simulations in which realizations of the predictors are generated as $X \sim U(0, 1)$.

To assess finite-sample marginal validity we calculate the proportion of responses that fall within the prediction region. To assess finite-sample local validity with respect to binning we first bin the predictor data and then, for each bin, we calculate the proportion of responses that fall within the prediction region. The same procedure is used to assess finite-sample conditional validity, but we use a much finer binning regime than what was used to assess finite-sample local validity.

The following function computes our estimate of prediction error at the observed data when the prediction region is not the nonparametric conformal prediction region.

```
absolute.error <- function(y, region, squared = TRUE){
  lwr <- region[, 1]
  upr <- region[, 2]
  index <- which(!(lwr <= y & y <= upr))
  out <- sum(unlist(lapply(index, function(j){
    error <- NULL
    if(squared == FALSE) error <-
      min(abs(y[j] - lwr[j]), abs(y[j] - upr[j]))
    if(squared == TRUE) error <-
      (min(abs(y[j] - lwr[j]), abs(y[j] - upr[j])))^2
    error
  }))) / length(y)
  out
}
```

The following function computes our estimate of prediction error at the observed data when the prediction region is the nonparametric conformal prediction region.

```

absolute.error.nonparametric <- function(data, region,
  squared = TRUE){
  n <- nrow(data)
  n.bins.region <- length(region)
  d <- ncol(data) - 1
  x <- as.matrix(data[,1:d + 1], col = d)
  index.bins.region <- find.index(x, wn = 1/n.bins.region, d = d)
  y <- data[, 1]
  area <- 0
  for(j in 1:n){
    index.j <- which(index.bins.region == j)
    error <- c(y[j] - region[[index.bins.region[j]]])
    index <- 0
    if(any(error > 0)) index <- max(which(error > 0))
    if(index %% 2 == 0){
      if(squared == FALSE) area <-
        area + min(abs(error)) / n
      if(squared == TRUE) area <-
        area + min(abs(error))^2 / n
    }
  }
  area
}

```

Our Monte Carlo simulator follows. This function computes all of the diagnostics and local coverage probabilities for each prediction region for every generated dataset.

```

gamma.simulator <- function(n = 500, alpha = 0.10, beta, bins = 3,
  family = "Gamma", link = "inverse", shape = 2,
  parametric = TRUE, nonparametric = TRUE,
  LS = TRUE, LSLW = TRUE, HD = TRUE, cores = 6){

  ## in this univariate problem, p and d are the same
  p <- d <- length(beta) - 1
  x <- matrix(runif(n), ncol = p)
  y <- rep(0, n)
  data <- NULL

  ## set up partition
  if(class(bins) == "NULL"){
    wn <- min(1/ floor(1 / (log(n)/n)^(1/(d+3))), 1/2)
    bins <- 1 / wn
  }

  ## generate the data (has functionality for different
  ## families and link functions)
  if(family == "Gamma"){

```

```

if(link == "identity"){
  rate <- (1 / (cbind(1, x) %*% beta)) * shape
  y <- rgamma(n = n, shape = shape, rate = rate)
  y <- y / sd(y)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
}
if(link == "inverse"){
  rate <- (cbind(1, x) %*% beta) * shape
  y <- rgamma(n = n, shape = shape, rate = rate)
  y <- y / sd(y)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
}
if(link == "log"){
  rate <- (1 / exp(cbind(1, x) %*% beta)) * shape
  y <- rgamma(n = n, shape = shape, rate = rate)
  y <- y / sd(y)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
}
}

if(family == "gaussian"){
  mu <- cbind(1, x) %*% beta
  y <- rnorm(n = n, mean = mu, sd = sd)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
}

if(family == "inverse.gaussian"){
  mu = 1 / sqrt(cbind(1, x) %*% beta)
  y <- rinvgauss(n = n, mean = mu)
  y <- y / sd(y)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
}

## fit the Gamma regression model
fit <- glm(y ~ x1, family = "Gamma", data = data)
paraCI <- nonparaCI <- LSCI <- LSLWCI <- HDCI <- NULL
formula <- fit$formula
newdata <- data
respname <- all.vars(formula)[1]
newdata <- newdata[, !(colnames(data) %in% respname)]
newdata <- as.matrix(newdata)

```

```

## obtain the prediction regions
if(parametric){
  cpred <- conformal.glm(fit, parametric = TRUE,
    nonparametric = FALSE, alpha = alpha,
    bins = bins, cores = cores)
  paraCI <- cpred$paraconformal
}
if(nonparametric){
  cpred <- conformal.glm(fit, parametric = FALSE,
    nonparametric = TRUE, alpha = alpha,
    bins = bins, cores = cores)
  nonparaCI <- cpred$nonparaconformal
}
if(LS){
  p1.tibs <- conformal.pred(x = cbind(x,x^2,x^3), y = y,
    x0 = cbind(x,x^2,x^3),
    train.fun = train.fun, predict.fun = predict.fun,
    alpha = alpha)
  LSCI <- cbind(p1.tibs$lo, p1.tibs$up)
}
if(LSLW){
  cubic.model <- lm(y ~ x + I(x^2) + I(x^3))
  abs.resid <- abs(cubic.model$resid)
  smooth.call <- smooth.spline(x, abs.resid,
    nknots = 10)
  lambda <- smooth.call$lambda
  df <- smooth.call$df
  mad.train.fun <- function(x, y, out = NULL){
    smooth.spline(x[, 1], y, lambda = lambda,
      df = df, nknots = 10)
  }
  p2.tibs <- conformal.pred(x = cbind(x,x^2,x^3), y = y,
    x0 = cbind(x,x^2,x^3),
    train.fun = train.fun, predict.fun = predict.fun,
    mad.train.fun = mad.train.fun,
    mad.predict.fun = mad.predict.fun,
    alpha = alpha)
  LSLWCI <- cbind(p2.tibs$lo, p2.tibs$up)
}
if(HD){
  betaMLE <- coefficients(fit)
  shapeMLE <- as.numeric(gamma.shape(fit)[1])
  rateMLE <- cbind(1, newdata) %*% betaMLE * shapeMLE
  HDCI <- do.call(rbind, lapply(1:nrow(newdata), function(j){
    hdi(qgamma, 1 - alpha, shape = shapeMLE, rate = rateMLE[j, 1])
  })))
}

```

```

## local coverage prediction regions
output.parametric <- output.nonparametric <-
  output.LS <- output.LSLW <- output.HD <- rep(NA, bins + 1)
if(parametric){
  marginal.parametric <- local.coverage(region = paraCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.parametric <- local.coverage(region = paraCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.parametric <- local.coverage(region = paraCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.parametric <- c(marginal.parametric, local.parametric,
    local.inx.parametric,
    mean(apply(paraCI, 1, diff)),
    absolute.error(y = y, region = paraCI))
}
if(nonparametric){
  marginal.nonparametric <- local.coverage(region = nonparaCI,
    nonparametric = "TRUE", data = data, d = p, bins = 1,
    at.data = "TRUE")
  local.nonparametric <- local.coverage(region = nonparaCI,
    nonparametric = "TRUE", data = data, d = p, bins = bins,
    at.data = "TRUE")
  local.inx.nonparametric <- local.coverage(region = nonparaCI,
    nonparametric = "TRUE", data = data, d = p, bins = 25,
    at.data = "TRUE")
  output.nonparametric <-
    c(marginal.nonparametric, local.nonparametric,
      local.inx.nonparametric,
      area.nonparametric(nonparaCI),
      absolute.error.nonparametric(data = data,
        region = nonparaCI))
}
if(LS){
  marginal.LS <- local.coverage(region = LSCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.LS <- local.coverage(region = LSCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.LS <- local.coverage(region = LSCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.LS <- c(marginal.LS, local.LS, local.inx.LS,
    mean(apply(LSCI, 1, diff)),
    absolute.error(y = y, region = LSCI))
}
if(LSLW){
  marginal.LSLW <- local.coverage(region = LSLWCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.LSLW <- local.coverage(region = LSLWCI,

```

```

    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.LSLW <- local.coverage(region = LSLWCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.LSLW <- c(marginal.LSLW, local.LSLW, local.inx.LSLW,
    mean(apply(LSLWCI, 1, diff)),
    absolute.error(y = y, region = LSLWCI))
}
if(HD){
  marginal.HD <- local.coverage(region = HDCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.HD <- local.coverage(region = HDCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.HD <- local.coverage(region = HDCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.HD <- c(marginal.HD, local.HD, local.inx.HD,
    mean(apply(HDCI, 1, diff)),
    absolute.error(y = y, region = HDCI))
}

output <- list(output.parametric = output.parametric,
  output.nonparametric = output.nonparametric,
  output.LS = output.LS,
  output.LSLW = output.LSLW,
  output.HD = output.HD)
output
}

```

The following performs our Monte Carlo simulation with $B = 50$ iterations.

```

B <- 50
system.time(local.500.3.2 <- do.call(cbind, lapply(1:B,
  FUN = function(j){
    unlist(gamma.simulator(beta = beta))
  })))

##      user      system    elapsed
## 13157.579    25.694    6385.948

```

```

local.gamma.500.3.2 <- cbind(
  rowMeans(local.500.3.2, na.rm = TRUE),
  apply(local.500.3.2, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

	parametric conformal	nonparametric conformal	LS conformal	LSLW conformal	HD region
marginal coverage	0.908 (0.0003)	0.903 (0.0002)	0.904 (0.0004)	0.915 (0.0008)	0.9 (0.0012)
local coverage when $0 < x < 1/3$	0.909 (0.0007)	0.903 (0.0006)	0.75 (0.0024)	0.913 (0.0016)	0.902 (0.0026)
local coverage when $1/3 \leq x < 2/3$	0.908 (0.0004)	0.902 (0.0003)	0.967 (0.0015)	0.908 (0.0024)	0.904 (0.003)
local coverage when $2/3 \leq x < 1$	0.908 (0.0004)	0.903 (0.0003)	0.994 (0.0008)	0.925 (0.0022)	0.895 (0.0029)
area	1.825 (0.02)	2.057 (0.0229)	2.342 (0.0175)	1.963 (0.0152)	1.792 (0.0167)
prediction error	0.083 (0.0055)	0.136 (0.0076)	0.129 (0.0057)	0.054 (0.0034)	0.07 (0.0038)

Table 1: Diagnostics of prediction regions. This table gives the area, prediction error, marginal coverage, and local coverages with respect to our binning scheme for the parametric, nonparametric, LS, and LSLW conformal prediction regions and the HD prediction region.

```
options(scipen = 999)
local.gamma.500.3.2[, 1] <- round(local.gamma.500.3.2[, 1], 3)
local.gamma.500.3.2[, 2] <- round(local.gamma.500.3.2[, 2], 4)
```

Diagnostics for each of the five prediction regions are given in Table 1 and Figure 2. We see that the parametric conformal prediction region performs as advertised. When the model is correctly specified, the parametric conformal prediction region is similar to the minimum length HD prediction region in area and prediction error (and in appearance as seen in Figure 1). The parametric conformal prediction region exhibits some conservative overcoverage marginally and with respect to binning, and some undercoverage in x for values close to the break points of the bins. The LSLW conformal prediction region has lower prediction error than the parametric conformal prediction error but such a benefit comes with the cost of lack of precision (increase in size) and dramatic overcoverage. The low prediction error of the LSLW conformal prediction region appears to stem from its ability to be far wider than the other prediction intervals at the values of x where the response data is most variable, as seen in Figure 1. This feature combined with its symmetric errors construction is what leads to its increase in size. The LSLW conformal prediction region is seen to provide conservative finite-sample local coverage in x . The nonparametric and LS conformal prediction regions are larger and have larger prediction error than the parametric conformal prediction region. The LS conformal prediction exhibits large undercoverage when the predictor is small and large overcoverage when the predictor is large. This is best evidenced by Figure 2.




Figure 2: Plot of the estimated coverage probabilities of prediction regions across x .

4 Additional Gamma simulations

For the simulations in this section we generate responses using a gamma regression model with one variable. We specify the inverse link function (the default in the `glm` function) and we set $\beta = (1.25, -1)^T$. This value of β is chosen so that the generated Gamma data is increasing, on average, in x and exhibits increasing variability in x . We investigate the performance of the five prediction regions across different sample size and shape parameter combinations. We consider sample sizes of $n \in \{150, 250, 500\}$ and shape parameter values of $\{0.5, 0.75, 1, 1.5, 2, 5, 7, 10, 12, 15, 20, 25, 30, 40, 50, 60, 70, 80, 90, 100\}$. The LS and LSLW conformal prediction regions are fit using a cubic regression model. This model is simple and it fits this type of Gamma data better than a simple linear or quadratic regression model. Note that as the shape parameter increases the cubic regression model fits the data better. When $n = 150$ we build the parametric and nonparametric conformal prediction regions using 2 bins. When $n = 250, 500$ we build the parametric and nonparametric conformal prediction regions using 3 bins. These number of bin choices correspond to the bin width asymptotics of Lei and Wasserman [2014].

The next subsection contains all the code necessary to reproduce our results in Section 4.2.

4.1 Simulations

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 0.5.

```
set.seed(13)
beta <- c(1.25, -1)
bins <- 2
```



```

n <- 150
B <- 250
system.time(out.gamma.150.2.0.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 0.5))
  })))

##      user      system elapsed
## 7473.082      70.418  4733.124

```

```

gamma.150.2.0.5 <- cbind(
  rowMeans(out.gamma.150.2.0.5, na.rm = TRUE),
  apply(out.gamma.150.2.0.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 0.75.

```

system.time(out.gamma.150.2.0.75 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 0.75))
  })))

##      user      system elapsed
## 7308.074      67.704  4385.919

```

```

gamma.150.2.0.75 <- cbind(
  rowMeans(out.gamma.150.2.0.75, na.rm = TRUE),
  apply(out.gamma.150.2.0.75, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 1.

```

system.time(out.gamma.150.2.1 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 1))
  })))

##      user      system elapsed
## 6992.204    67.186 4018.650

```

```

gamma.150.2.1 <- cbind(
  rowMeans(out.gamma.150.2.1, na.rm = TRUE),
  apply(out.gamma.150.2.1, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 1.5.

```

system.time(out.gamma.150.2.1.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 1.5))
  })))

##      user      system elapsed
## 7456.522    67.902 10638.787

```

```

gamma.150.2.1.5 <- cbind(
  rowMeans(out.gamma.150.2.1.5, na.rm = TRUE),
  apply(out.gamma.150.2.1.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 2.

```

system.time(out.gamma.150.2.2 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 2))
  })))

```

```
##      user      system elapsed
## 8411.829      71.428 4583.947
```

```
gamma.150.2.2 <- cbind(
  rowMeans(out.gamma.150.2.2, na.rm = TRUE),
  apply(out.gamma.150.2.2, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 5.

```
system.time(out.gamma.150.2.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 5))
  })))
```

```
##      user      system elapsed
## 7584.710      64.449 4059.795
```

```
gamma.150.2.5 <- cbind(
  rowMeans(out.gamma.150.2.5, na.rm = TRUE),
  apply(out.gamma.150.2.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 7.

```
system.time(out.gamma.150.2.7 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 7))
  })))
```

```
##      user      system elapsed
## 7562.858      65.481 4018.408
```

```
gamma.150.2.7 <- cbind(
  rowMeans(out.gamma.150.2.7, na.rm = TRUE),
  apply(out.gamma.150.2.7, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 10.

```
system.time(out.gamma.150.2.10 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 10))
  })))

##      user      system elapsed
## 7645.495      66.652 4035.760
```

```
gamma.150.2.10 <- cbind(
  rowMeans(out.gamma.150.2.10, na.rm = TRUE),
  apply(out.gamma.150.2.10, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 12.

```
system.time(out.gamma.150.2.12 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 12))
  })))

##      user      system elapsed
## 7641.583      65.990 4035.577
```

```
gamma.150.2.12 <- cbind(
  rowMeans(out.gamma.150.2.12, na.rm = TRUE),
  apply(out.gamma.150.2.12, 1,
    FUN = function(x){
```

```

sds <- sd(x, na.rm = TRUE)
lengths <- length(which(!is.na(x)))
sds / sqrt(lengths)
}))

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 15.

```

system.time(out.gamma.150.2.15 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 15))
  })))

##      user      system elapsed
## 7535.949    65.387  4003.467

```

```

gamma.150.2.15 <- cbind(
  rowMeans(out.gamma.150.2.15, na.rm = TRUE),
  apply(out.gamma.150.2.15, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 20.

```

system.time(out.gamma.150.2.20 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 20))
  })))

##      user      system elapsed
## 7850.360    66.055  4065.622

```

```

gamma.150.2.20 <- cbind(
  rowMeans(out.gamma.150.2.20, na.rm = TRUE),
  apply(out.gamma.150.2.20, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 25.

```
system.time(out.gamma.150.2.25 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 25))
  })))

##      user      system elapsed
## 7805.004    66.341 4057.034
```

```
gamma.150.2.25 <- cbind(
  rowMeans(out.gamma.150.2.25, na.rm = TRUE),
  apply(out.gamma.150.2.25, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 30.

```
system.time(out.gamma.150.2.30 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 30))
  })))

##      user      system elapsed
## 8044.189    68.723 4100.474
```

```
gamma.150.2.30 <- cbind(
  rowMeans(out.gamma.150.2.30, na.rm = TRUE),
  apply(out.gamma.150.2.30, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 40.

```

system.time(out.gamma.150.2.40 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 40))
  })))

##      user      system elapsed
## 7968.318    65.667 4073.458

```

```

gamma.150.2.40 <- cbind(
  rowMeans(out.gamma.150.2.40, na.rm = TRUE),
  apply(out.gamma.150.2.40, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 50.

```

system.time(out.gamma.150.2.50 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 50))
  })))

##      user      system elapsed
## 8087.560    66.821 4086.898

```

```

gamma.150.2.50 <- cbind(
  rowMeans(out.gamma.150.2.50, na.rm = TRUE),
  apply(out.gamma.150.2.50, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 60.

```

system.time(out.gamma.150.2.60 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 60))
  })))

```

```
##      user      system elapsed
## 8543.967      70.317 4309.284
```

```
gamma.150.2.60 <- cbind(
  rowMeans(out.gamma.150.2.60, na.rm = TRUE),
  apply(out.gamma.150.2.60, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 70.

```
system.time(out.gamma.150.2.70 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 70))
  })))
```

```
##      user      system elapsed
## 9038.847      73.108 4394.150
```

```
gamma.150.2.70 <- cbind(
  rowMeans(out.gamma.150.2.70, na.rm = TRUE),
  apply(out.gamma.150.2.70, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 80.

```
system.time(out.gamma.150.2.80 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 80))
  })))
```

```
##      user      system elapsed
## 9580.813      75.711 4638.776
```



```
gamma.150.2.80 <- cbind(
  rowMeans(out.gamma.150.2.80, na.rm = TRUE),
  apply(out.gamma.150.2.80, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 90.

```
system.time(out.gamma.150.2.90 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 90))
  })))

##      user      system elapsed
## 9410.124    74.038  4608.842
```

```
gamma.150.2.90 <- cbind(
  rowMeans(out.gamma.150.2.90, na.rm = TRUE),
  apply(out.gamma.150.2.90, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 100.

```
system.time(out.gamma.150.2.100 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 100))
  })))

##      user      system elapsed
## 9667.834    76.044  4650.667
```

```
gamma.150.2.100 <- cbind(
  rowMeans(out.gamma.150.2.100, na.rm = TRUE),
  apply(out.gamma.150.2.100, 1,
    FUN = function(x){
```

```

sds <- sd(x, na.rm = TRUE)
lengths <- length(which(!is.na(x)))
sds / sqrt(lengths)
}))

```

Reorganize the output.

```

para.area.gamma.150 <- nonpara.area.gamma.150 <-
  LS.area.gamma.150 <- LSLW.area.gamma.150 <-
  HD.area.gamma.150 <- NULL
para.error.gamma.150 <- nonpara.error.gamma.150 <-
  LS.error.gamma.150 <- LSLW.error.gamma.150 <-
  HD.error.gamma.150 <- NULL
para.marginal.gamma.150 <- nonpara.marginal.gamma.150 <-
  LS.marginal.gamma.150 <- LSLW.marginal.gamma.150 <-
  HD.marginal.gamma.150 <- NULL
para.local.gamma.150 <- nonpara.local.gamma.150 <-
  LS.local.gamma.150 <- LSLW.local.gamma.150 <-
  HD.local.gamma.150 <- NULL
para.inx.gamma.150 <- nonpara.inx.gamma.150 <-
  LS.inx.gamma.150 <- LSLW.inx.gamma.150 <-
  HD.inx.gamma.150 <- NULL
shapes <- c(0.5, 0.75, 1, 1.5, 2, 5, 7, 10, 12, 15, 20, 25,
  30, 40, 50, 60, 70, 80, 90, 100)
for(j in shapes ){
  internal.output <-
    eval(parse(text=paste("gamma.150.2", j, sep = ".")))
  k <- which(shapes == j)
  para.marginal.gamma.150[k] <- as.numeric(internal.output[1, 1])
  para.local.gamma.150[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[2:3, 1])
  para.inx.gamma.150[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[4:28, 1])
  para.error.gamma.150[k] <- as.numeric(internal.output[29, 1])
  para.area.gamma.150[k] <- as.numeric(internal.output[30, 1])
  nonpara.marginal.gamma.150[k] <- as.numeric(internal.output[31, 1])
  nonpara.local.gamma.150[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[32:33, 1])
  nonpara.inx.gamma.150[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[34:58, 1])
  nonpara.error.gamma.150[k] <- as.numeric(internal.output[59, 1])
  nonpara.area.gamma.150[k] <- as.numeric(internal.output[60, 1])
  LS.marginal.gamma.150[k] <- as.numeric(internal.output[61, 1])
  LS.local.gamma.150[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[62:63, 1])
  LS.inx.gamma.150[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[64:88, 1])
  LS.error.gamma.150[k] <- as.numeric(internal.output[89, 1])

```

```

LS.area.gamma.150[k] <- as.numeric(internal.output[90, 1])
LSLW.marginal.gamma.150[k] <- as.numeric(internal.output[91, 1])
LSLW.local.gamma.150[c((bins*(k-1)+1):(bins*k))] <-
  as.numeric(internal.output[92:93, 1])
LSLW.inx.gamma.150[c((25*(k-1)+1):(25*k))] <-
  as.numeric(internal.output[94:118, 1])
LSLW.error.gamma.150[k] <- as.numeric(internal.output[119, 1])
LSLW.area.gamma.150[k] <- as.numeric(internal.output[120, 1])
HD.marginal.gamma.150[k] <- as.numeric(internal.output[121, 1])
HD.local.gamma.150[c((bins*(k-1)+1):(bins*k))] <-
  as.numeric(internal.output[122:123, 1])
HD.inx.gamma.150[c((25*(k-1)+1):(25*k))] <-
  as.numeric(internal.output[124:148, 1])
HD.error.gamma.150[k] <- as.numeric(internal.output[149, 1])
HD.area.gamma.150[k] <- as.numeric(internal.output[150, 1])
}

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 0.5.

```

set.seed(13)
beta <- c(1.25, -1)
bins <- 3
n <- 250
B <- 50
system.time(out.gamma.250.3.0.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 0.5))
  })))

##      user      system elapsed
## 2511.339      16.860 1496.168

```

```

gamma.250.3.0.5 <- cbind(
  rowMeans(out.gamma.250.3.0.5, na.rm = TRUE),
  apply(out.gamma.250.3.0.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 0.75.

```

system.time(out.gamma.250.3.0.75 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 0.75))
  })))

##      user      system elapsed
## 2577.735    16.795 1500.989

```

```

gamma.250.3.0.75 <- cbind(
  rowMeans(out.gamma.250.3.0.75, na.rm = TRUE),
  apply(out.gamma.250.3.0.75, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and $\text{shape} = 1$.

```

system.time(out.gamma.250.3.1 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 1))
  })))

##      user      system elapsed
## 2694.797    16.807 1523.346

```

```

gamma.250.3.1 <- cbind(
  rowMeans(out.gamma.250.3.1, na.rm = TRUE),
  apply(out.gamma.250.3.1, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and $\text{shape} = 1.5$.

```

system.time(out.gamma.250.3.1.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 1.5))
  })))

```

```
##      user      system elapsed
## 2820.080      17.131 1561.781
```

```
gamma.250.3.1.5 <- cbind(
  rowMeans(out.gamma.250.3.1.5, na.rm = TRUE),
  apply(out.gamma.250.3.1.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 2.

```
system.time(out.gamma.250.3.2 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 2))
  })))

##      user      system elapsed
## 3105.011      16.995 1603.930
```

```
gamma.250.3.2 <- cbind(
  rowMeans(out.gamma.250.3.2, na.rm = TRUE),
  apply(out.gamma.250.3.2, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 5.

```
system.time(out.gamma.250.3.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 5))
  })))

##      user      system elapsed
## 2969.224      17.372 1572.685
```

```
gamma.250.3.5 <- cbind(
  rowMeans(out.gamma.250.3.5, na.rm = TRUE),
  apply(out.gamma.250.3.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 7.

```
system.time(out.gamma.250.3.7 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 7))
  })))

##      user      system elapsed
## 2935.660      17.456 1564.401
```

```
gamma.250.3.7 <- cbind(
  rowMeans(out.gamma.250.3.7, na.rm = TRUE),
  apply(out.gamma.250.3.7, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 10.

```
system.time(out.gamma.250.3.10 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 10))
  })))

##      user      system elapsed
## 2842.656      17.363 1539.173
```

```
gamma.250.3.10 <- cbind(
  rowMeans(out.gamma.250.3.10, na.rm = TRUE),
  apply(out.gamma.250.3.10, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
```

```
lengths <- length(which(!is.na(x)))
sds / sqrt(lengths)
}))
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 12.

```
system.time(out.gamma.250.3.12 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 12))
  })))

##      user      system elapsed
## 2760.976    17.350 1526.356
```

```
gamma.250.3.12 <- cbind(
  rowMeans(out.gamma.250.3.12, na.rm = TRUE),
  apply(out.gamma.250.3.12, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 15.

```
system.time(out.gamma.250.3.15 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 15))
  })))

##      user      system elapsed
## 2729.205    17.260 1519.564
```

```
gamma.250.3.15 <- cbind(
  rowMeans(out.gamma.250.3.15, na.rm = TRUE),
  apply(out.gamma.250.3.15, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 20.

```
system.time(out.gamma.250.3.20 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 20))
  })))

##      user      system elapsed
## 2701.140    17.335 1504.800
```

```
gamma.250.3.20 <- cbind(
  rowMeans(out.gamma.250.3.20, na.rm = TRUE),
  apply(out.gamma.250.3.20, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 25.

```
system.time(out.gamma.250.3.25 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 25))
  })))

##      user      system elapsed
## 2881.666    17.454 1544.467
```

```
gamma.250.3.25 <- cbind(
  rowMeans(out.gamma.250.3.25, na.rm = TRUE),
  apply(out.gamma.250.3.25, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

Reorganize the output.

```
para.area.gamma.250 <- nonpara.area.gamma.250 <-
LS.area.gamma.250 <- LSLW.area.gamma.250 <-
```



```

HD.area.gamma.250 <- NULL
para.error.gamma.250 <- nonpara.error.gamma.250 <-
  LS.error.gamma.250 <- LSLW.error.gamma.250 <-
  HD.error.gamma.250 <- NULL
para.marginal.gamma.250 <- nonpara.marginal.gamma.250 <-
  LS.marginal.gamma.250 <- LSLW.marginal.gamma.250 <-
  HD.marginal.gamma.250 <- NULL
para.local.gamma.250 <- nonpara.local.gamma.250 <-
  LS.local.gamma.250 <- LSLW.local.gamma.250 <-
  HD.local.gamma.250 <- NULL
para.inx.gamma.250 <- nonpara.inx.gamma.250 <-
  LS.inx.gamma.250 <- LSLW.inx.gamma.250 <-
  HD.inx.gamma.250 <- NULL
shapes <- c(0.5, 0.75, 1, 1.5, 2, 5, 7, 10, 12, 15, 20, 25)
for(j in shapes ){
  internal.output <- eval(parse(text=paste("gamma.250.3", j, sep = ".")))
  k <- which(shapes == j)
  para.marginal.gamma.250[k] <- as.numeric(internal.output[1, 1])
  para.local.gamma.250[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[2:4, 1])
  para.inx.gamma.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[5:29, 1])
  para.error.gamma.250[k] <- as.numeric(internal.output[30, 1])
  para.area.gamma.250[k] <- as.numeric(internal.output[31, 1])
  nonpara.marginal.gamma.250[k] <- as.numeric(internal.output[32, 1])
  nonpara.local.gamma.250[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[33:35, 1])
  nonpara.inx.gamma.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[36:60, 1])
  nonpara.error.gamma.250[k] <- as.numeric(internal.output[61, 1])
  nonpara.area.gamma.250[k] <- as.numeric(internal.output[62, 1])
  LS.marginal.gamma.250[k] <- as.numeric(internal.output[63, 1])
  LS.local.gamma.250[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[64:66, 1])
  LS.inx.gamma.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[67:91, 1])
  LS.error.gamma.250[k] <- as.numeric(internal.output[92, 1])
  LS.area.gamma.250[k] <- as.numeric(internal.output[93, 1])
  LSLW.marginal.gamma.250[k] <- as.numeric(internal.output[94, 1])
  LSLW.local.gamma.250[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[95:97, 1])
  LSLW.inx.gamma.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[98:122, 1])
  LSLW.error.gamma.250[k] <- as.numeric(internal.output[123, 1])
  LSLW.area.gamma.250[k] <- as.numeric(internal.output[124, 1])
  HD.marginal.gamma.250[k] <- as.numeric(internal.output[125, 1])
  HD.local.gamma.250[c((bins*(k-1)+1):(bins*k))] <-

```

```

    as.numeric(internal.output[126:128, 1])
HD.inx.gamma.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[129:153, 1])
HD.error.gamma.250[k] <- as.numeric(internal.output[154, 1])
HD.area.gamma.250[k] <- as.numeric(internal.output[155, 1])
}

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 0.5.

```

set.seed(13)
beta <- c(1.25, -1)
bins <- 3
n <- 500
B <- 50
system.time(out.gamma.500.3.0.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 0.5))
  })))

##      user      system elapsed
## 9094.705    23.855 5496.216

```

```

gamma.500.3.0.5 <- cbind(
  rowMeans(out.gamma.500.3.0.5, na.rm = TRUE),
  apply(out.gamma.500.3.0.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 0.75.

```

system.time(out.gamma.500.3.0.75 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 0.75))
  })))

##      user      system elapsed
## 8411.859    21.506 4823.067

```

```
gamma.500.3.0.75 <- cbind(
  rowMeans(out.gamma.500.3.0.75, na.rm = TRUE),
  apply(out.gamma.500.3.0.75, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 1.

```
system.time(out.gamma.500.3.1 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 1))
  })))

##      user      system elapsed
## 9375.044      23.025 5350.013
```

```
gamma.500.3.1 <- cbind(
  rowMeans(out.gamma.500.3.1, na.rm = TRUE),
  apply(out.gamma.500.3.1, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 1.5.

```
system.time(out.gamma.500.3.1.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 1.5))
  })))

##      user      system elapsed
## 9709.719      19.786 5023.810
```

```
gamma.500.3.1.5 <- cbind(
  rowMeans(out.gamma.500.3.1.5, na.rm = TRUE),
  apply(out.gamma.500.3.1.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
```

```
lengths <- length(which(!is.na(x)))
sds / sqrt(lengths)
}))
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and $\text{shape} = 2$.

```
system.time(out.gamma.500.3.2 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 2))
  })))
```

```
##      user      system    elapsed
## 11685.253      20.810    5403.206
```

```
gamma.500.3.2 <- cbind(
  rowMeans(out.gamma.500.3.2, na.rm = TRUE),
  apply(out.gamma.500.3.2, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and $\text{shape} = 5$.

```
system.time(out.gamma.500.3.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 5))
  })))
```

```
##      user      system    elapsed
## 10175.065      20.472    5060.978
```

```
gamma.500.3.5 <- cbind(
  rowMeans(out.gamma.500.3.5, na.rm = TRUE),
  apply(out.gamma.500.3.5, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and $\text{shape} = 7$.

```

system.time(out.gamma.500.3.7 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 7))
  })))

##          user      system    elapsed
## 10187.202      20.175 12619.584

```

```

gamma.500.3.7 <- cbind(
  rowMeans(out.gamma.500.3.7, na.rm = TRUE),
  apply(out.gamma.500.3.7, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 10.

```

system.time(out.gamma.500.3.10 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 10))
  })))

##          user      system    elapsed
## 10417.609      20.990  5142.265

```

```

gamma.500.3.10 <- cbind(
  rowMeans(out.gamma.500.3.10, na.rm = TRUE),
  apply(out.gamma.500.3.10, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 12.

```

system.time(out.gamma.500.3.12 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 12))
  })))

```

```
##      user      system elapsed
## 9750.862    19.680 4878.616
```

```
gamma.500.3.12 <- cbind(
  rowMeans(out.gamma.500.3.12, na.rm = TRUE),
  apply(out.gamma.500.3.12, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 15.

```
system.time(out.gamma.500.3.15 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 15))
  })))
```

```
##      user      system elapsed
## 9450.264    19.315 4806.449
```

```
gamma.500.3.15 <- cbind(
  rowMeans(out.gamma.500.3.15, na.rm = TRUE),
  apply(out.gamma.500.3.15, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 20.

```
system.time(out.gamma.500.3.20 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 20))
  })))
```

```
##      user      system elapsed
## 9359.494    18.956 4786.326
```

```
gamma.500.3.20 <- cbind(
  rowMeans(out.gamma.500.3.20, na.rm = TRUE),
  apply(out.gamma.500.3.20, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 25.

```
system.time(out.gamma.500.3.25 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(gamma.simulator(beta = beta, n = n,
      bins = bins, shape = 25))
  })))

##      user      system elapsed
## 9465.436    19.651 4784.069
```

```
gamma.500.3.25 <- cbind(
  rowMeans(out.gamma.500.3.25, na.rm = TRUE),
  apply(out.gamma.500.3.25, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

Reorganize the output.

```
para.area.gamma.500 <- nonpara.area.gamma.500 <-
  LS.area.gamma.500 <- LSLW.area.gamma.500 <-
  HD.area.gamma.500 <- NULL
para.error.gamma.500 <- nonpara.error.gamma.500 <-
  LS.error.gamma.500 <- LSLW.error.gamma.500 <-
  HD.error.gamma.500 <- NULL
para.marginal.gamma.500 <- nonpara.marginal.gamma.500 <-
  LS.marginal.gamma.500 <- LSLW.marginal.gamma.500 <-
  HD.marginal.gamma.500 <- NULL
para.local.gamma.500 <- nonpara.local.gamma.500 <-
  LS.local.gamma.500 <- LSLW.local.gamma.500 <-
  HD.local.gamma.500 <- NULL
para.inx.gamma.500 <- nonpara.inx.gamma.500 <-
  LS.inx.gamma.500 <- LSLW.inx.gamma.500 <-
  HD.inx.gamma.500 <- NULL
```

```

shapes <- c(0.5, 0.75, 1, 1.5, 2, 5, 7, 10, 12, 15, 20, 25)
for(j in shapes ){
  internal.output <- eval(parse(text=paste("gamma.500.3", j, sep = ".")))
  k <- which(shapes == j)
  para.marginal.gamma.500[k] <- as.numeric(internal.output[1, 1])
  para.local.gamma.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[2:4, 1])
  para.inx.gamma.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[5:29, 1])
  para.error.gamma.500[k] <- as.numeric(internal.output[30, 1])
  para.area.gamma.500[k] <- as.numeric(internal.output[31, 1])
  nonpara.marginal.gamma.500[k] <- as.numeric(internal.output[32, 1])
  nonpara.local.gamma.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[33:35, 1])
  nonpara.inx.gamma.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[36:60, 1])
  nonpara.error.gamma.500[k] <- as.numeric(internal.output[61, 1])
  nonpara.area.gamma.500[k] <- as.numeric(internal.output[62, 1])
  LSLW.marginal.gamma.500[k] <- as.numeric(internal.output[63, 1])
  LSLW.local.gamma.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[64:66, 1])
  LSLW.inx.gamma.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[67:91, 1])
  LSLW.error.gamma.500[k] <- as.numeric(internal.output[92, 1])
  LSLW.area.gamma.500[k] <- as.numeric(internal.output[93, 1])
  HD.marginal.gamma.500[k] <- as.numeric(internal.output[94, 1])
  HD.local.gamma.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[95:97, 1])
  HD.inx.gamma.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[98:122, 1])
  HD.error.gamma.500[k] <- as.numeric(internal.output[123, 1])
  HD.area.gamma.500[k] <- as.numeric(internal.output[124, 1])
  HD.marginal.gamma.500[k] <- as.numeric(internal.output[125, 1])
  HD.local.gamma.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[126:128, 1])
  HD.inx.gamma.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[129:153, 1])
  HD.error.gamma.500[k] <- as.numeric(internal.output[154, 1])
  HD.area.gamma.500[k] <- as.numeric(internal.output[155, 1])
}

```


4.2 Results

Results from our simulations are depicted in Figures 3-8. For all five prediction regions we depict the estimated area, prediction error, and local coverage probabilities with respect to binning in Figures 3, 5, and 7 for $n = 150, 250$, and 500 respectively. For all five prediction regions we depict the local coverage probabilities across x in Figures 4, 6, and 8 for $n = 150, 250$, and 500 respectively.

From these simulations we see that the parametric conformal prediction region is similar to the HD prediction region in area, prediction error, and appearance (as seen in Section 7.1). Moreover, the parametric conformal prediction region possess finite-sample, albeit slightly conservative, local validity with respect to binning and possess nominal finite-sample conditional coverage over most of the support. The LSLW conformal prediction region is smaller than both the parametric conformal prediction region and the HD prediction region. However, it has a higher prediction error than these regions and it is more conservative than the parametric conformal prediction region in most of our simulation settings. It also does not fit the data well when the deviations about the estimated mean function are clearly not symmetric as evidenced by Figures 16-19 in Section 7.1. These figures correspond to data generated from a Gamma distribution with small shape parameter values. The nonparametric conformal prediction region gives closer to nominal coverage and it possesses finite-sample local validity with respect to binning. However, this prediction region is larger and gives higher prediction error than the parametric and LSLW conformal prediction regions and HD prediction region. The LS conformal prediction region provides extreme overcoverage at small values of x and extreme undercoverage at large values of x . This prediction region is also larger and gives higher prediction error than the parametric and LSLW conformal prediction regions and HD prediction region.




Figure 3: This figure compares the performance of the parametric, nonparametric, least squares, and least squares locally weighted conformal prediction region and the highest density prediction region when $n = 150$ and the number of bins equals 2. The specific diagnostics used to compare these prediction regions is the area (top-left panel), prediction error (top-right panel), and the coverage probability with respect to binning (bottom row) across shape parameter values. The average of 250 Monte Carlo samples at each shape parameter value in these simulation settings form the lines that are depicted in this figure.

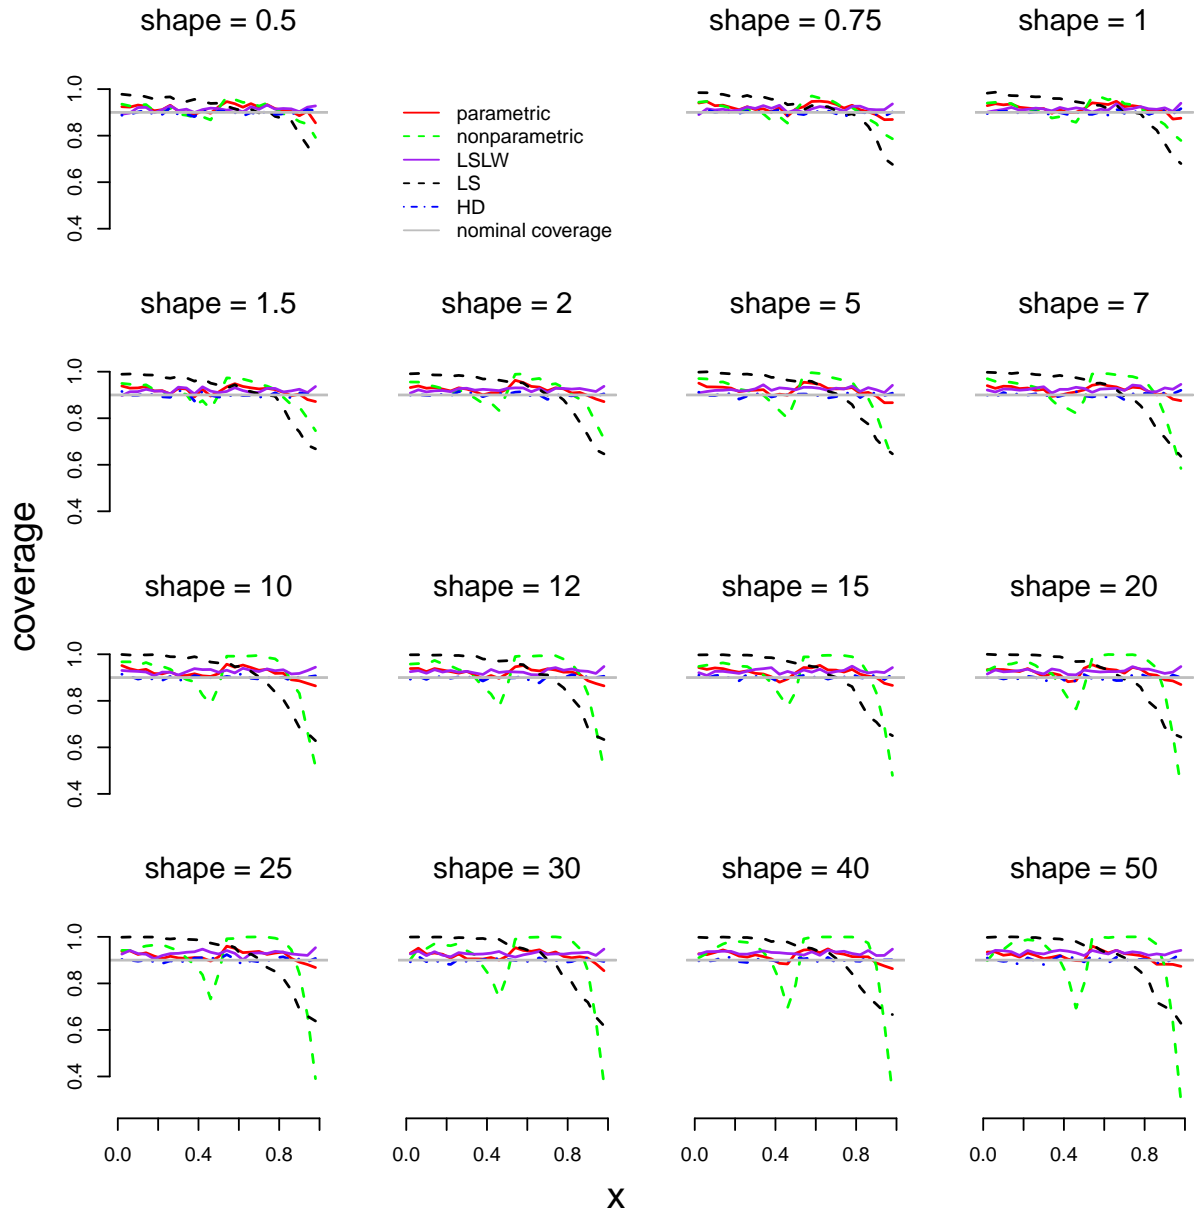


Figure 4: Plot of the estimated coverage probabilities of prediction regions across x and shape parameter values when the model is correctly specified, $n = 150$, and the number of bins is equal to 2.

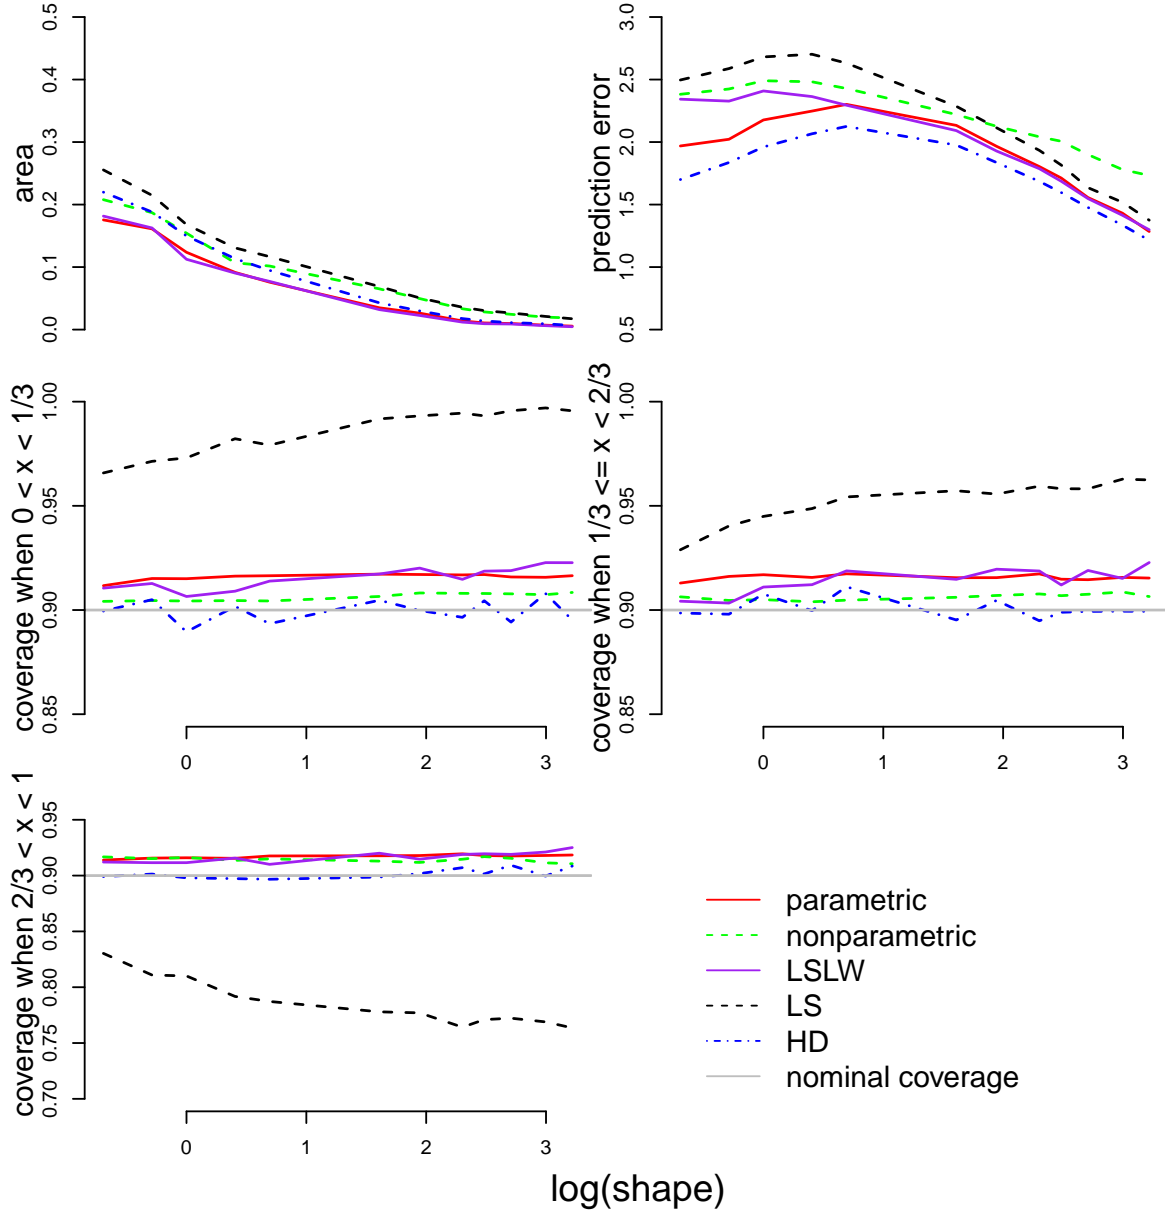


Figure 5: This figure compares the performance of the parametric, nonparametric, least squares, and least squares locally weighted conformal prediction region and the highest density prediction region when $n = 250$ and the number of bins equals 3. The specific diagnostics used to compare these prediction regions is the area (top-left panel), prediction error (top-right panel), and the coverage probability with respect to binning (bottom row) across shape parameter values. The average of 50 Monte Carlo samples at each shape parameter value in these simulation settings form the lines that are depicted in this figure.




Figure 6: Plot of the estimated coverage probabilities of prediction regions across x and shape parameter values when the model is correctly specified, $n = 250$, and the number of bins is equal to 3.

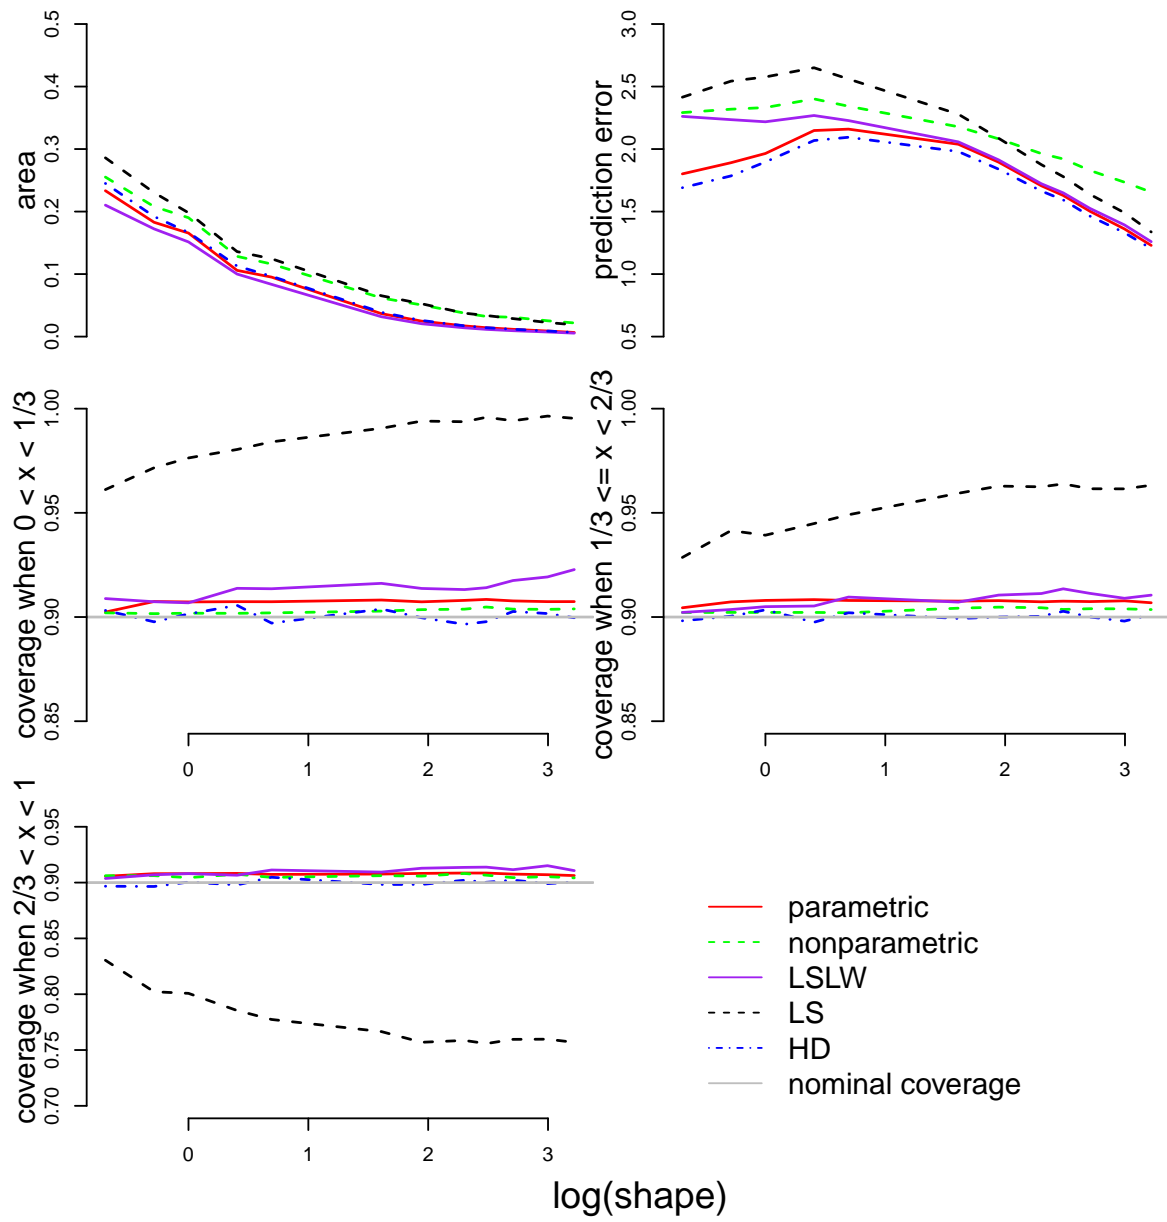


Figure 7: Diagnostic plots for prediction regions when the model is correctly specified, $n = 500$, and the number of bins is equal to 3.

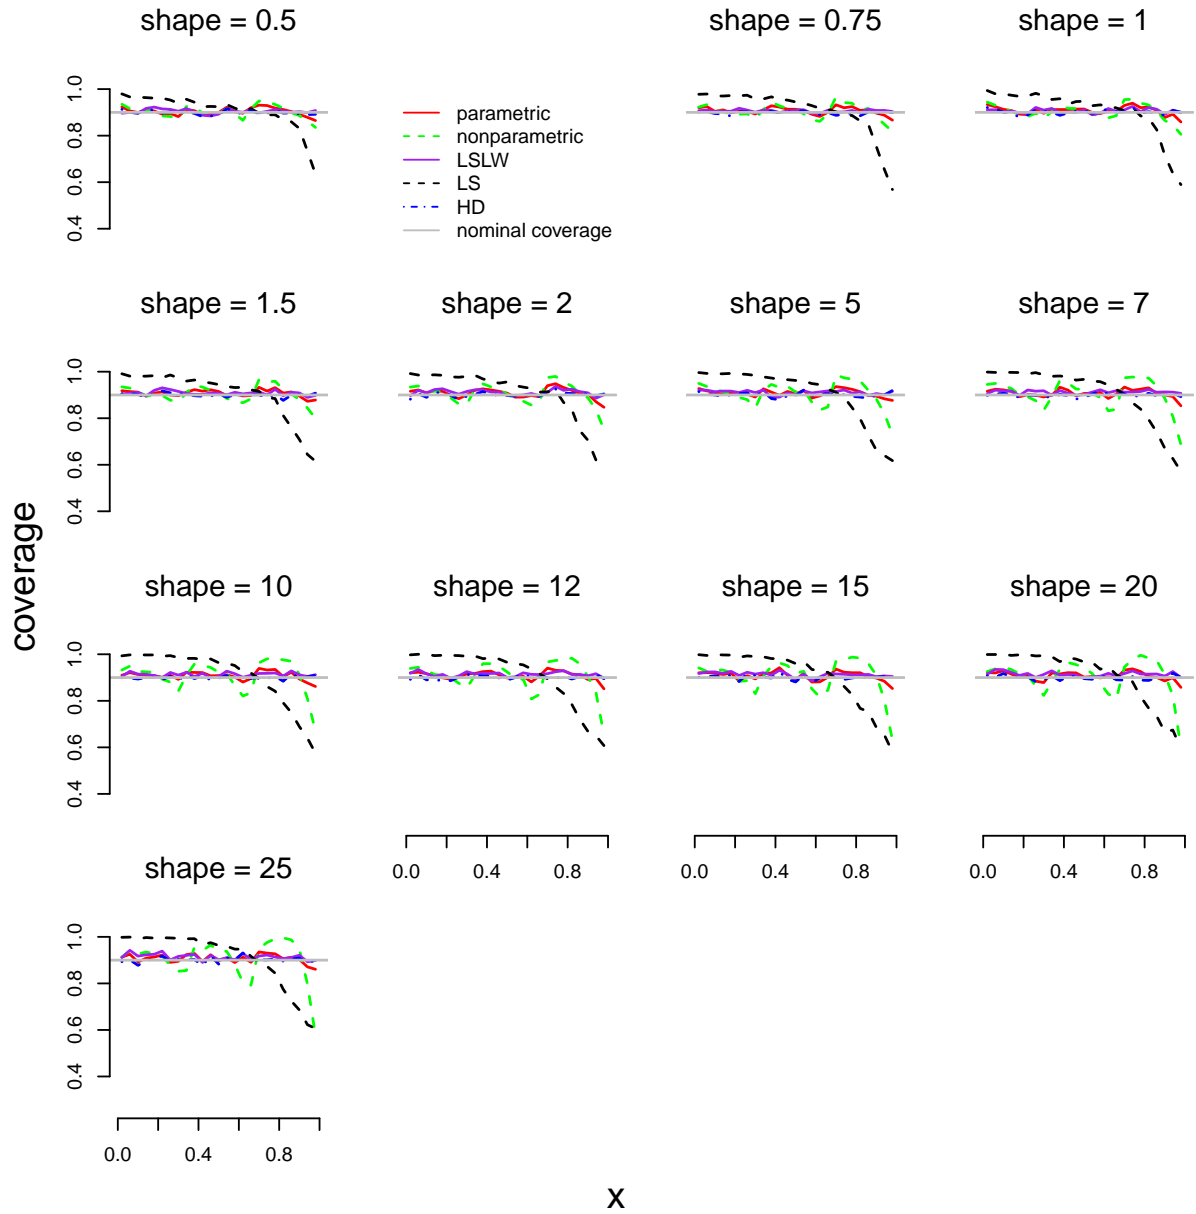


Figure 8: This figure compares the performance of the parametric, nonparametric, least squares, and least squares locally weighted conformal prediction region and the highest density prediction region when $n = 550$ and the number of bins equals 3. The specific diagnostics used to compare these prediction regions is the area (top-left panel), prediction error (top-right panel), and the coverage probability with respect to binning (bottom row) across shape parameter values. The average of 50 Monte Carlo samples at each shape parameter value in these simulation settings form the lines that are depicted in this figure.

5 Gamma-Gaussian model misspecification

In this Section, we compare the parametric conformal prediction region, the nonparametric conformal prediction region, the LSLW conformal prediction region, the LS conformal prediction region, and the HD prediction region under model misspecification. The data generating process is Gamma with an inverse link function and we set $\beta = (0.5, 1)^T$. We consider sample sizes of $n \in \{150, 250, 500\}$ and shape parameter values of $\{5, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100\}$ when $n = 150$ and shape parameter values of $\{5, 20, 40, 60, 80, 100\}$ when $n = 250, 500$. This value of β and these shape parameter values are chosen so that errors about a cubic regression model appear to be almost symmetric. In this analysis the parametric, LSLW, and LS conformal prediction regions and the HD prediction region are fit assuming this misspecified cubic regression model with homoscedastic normal errors. When $n = 150$ we build the parametric and nonparametric conformal prediction regions using 2 bins. When $n = 250, 500$ we build the parametric and nonparametric conformal prediction regions using 3 bins. These number of bin choices correspond to the bin width asymptotics of Lei and Wasserman [2014].

5.1 Simulations

The following function computes our diagnostic measures for the five prediction regions under investigation in the univariate case where data is assumed to be Gamma with inverse link function and the fitted model is Gaussian with a cubic fit.

```
misspec.simulator <- function(n = 500, alpha = 0.10, beta,
  shape, bins = 3, family = "Gamma", link = "inverse",
  confamily = "gaussian", parametric = TRUE,
  nonparametric = TRUE, LS = TRUE, LSLW = TRUE,
  HD = TRUE, cores = 6){

  p <- d <- length(beta) - 1
  x <- matrix(runif(n), ncol = p)
  y <- rep(0, n)
  data <- NULL

  ## set up partition
  if(class(bins) == "NULL"){
    wn <- min(1 / floor(1 / (log(n)/n)^(1/(d+3))), 1/2)
    bins <- 1 / wn
  }

  ## generate the data (has functionality for different
  ## families and link functions)
  if(family == "Gamma"){
    if(link == "identity"){
      rate <- (1 / (cbind(1, x) %*% beta)) * shape
      y <- rgamma(n = n, shape = shape, rate = rate)
      y <- y / sd(y)
      data <- data.frame(y = y, x = x)
      colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
    }
  }
```



```

if(link == "inverse"){
  rate <- (cbind(1, x) %*% beta) * shape
  y <- rgamma(n = n, shape = shape, rate = rate)
  y <- y / sd(y)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
}
if(link == "log"){
  rate <- (1 / exp(cbind(1, x) %*% beta)) * shape
  y <- rgamma(n = n, shape = shape, rate = rate)
  y <- y / sd(y)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
}
}

if(family == "gaussian"){
  mu <- cbind(1, x) %*% beta
  y <- rnorm(n = n, mean = mu, sd = sd)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
}

if(family == "inverse.gaussian"){
  mu = 1 / sqrt(cbind(1, x) %*% beta)
  y <- rinvgauss(n = n, mean = mu)
  y <- y / sd(y)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")
}

## fit the misspecified cubic regression model
fit <- glm(y ~ x1 + I(x1^2) + I(x1^3), family = confamily,
  data = data)
paraCI <- nonparaCI <- LSCI <- LSLWCI <- HDCI <- NULL
formula <- fit$formula
newdata <- data
respname <- all.vars(formula)[1]
newdata <- newdata[, !(colnames(data) %in% respname)]
newdata <- as.matrix(newdata)

## obtain the prediction regions
if(parametric){
  cpred <- conformal.glm(fit, parametric = TRUE,
    nonparametric = FALSE, alpha = alpha,
    bins = bins, cores = cores)
  paraCI <- cpred$paraconformal

```

```

}
if(nonparametric){
  cpred <- conformal.glm(fit, parametric = FALSE,
    nonparametric = TRUE, alpha = alpha,
    bins = bins, cores = cores)
  nonparaCI <- cpred$nonparaconformal
}
if(LS){
  p1.tibs <- conformal.pred(x = cbind(x, x^2, x^3), y = y,
    x0 = cbind(x, x^2, x^3),
    train.fun = train.fun, predict.fun = predict.fun,
    alpha = alpha)
  LSCI <- cbind(p1.tibs$lo, p1.tibs$up)
}
if(LSLW){
  cubic.model <- lm(y ~ x + I(x^2) + I(x^3))
  abs.resid <- abs(cubic.model$resid)
  smooth.call <- smooth.spline(x, abs.resid,
    nknots = 10)
  lambda <- smooth.call$lambda
  df <- smooth.call$df
  mad.train.fun <- function(x, y, out = NULL){
    smooth.spline(x[, 1], y, lambda = lambda,
      df = df, nknots = 10)
  }
  p2.tibs <- conformal.pred(x = cbind(x, x^2, x^3), y = y,
    x0 = cbind(x, x^2, x^3),
    train.fun = train.fun, predict.fun = predict.fun,
    mad.train.fun = mad.train.fun,
    mad.predict.fun = mad.predict.fun,
    alpha = alpha)
  LSLWCI <- cbind(p2.tibs$lo, p2.tibs$up)
}
if(HD){
  if(confamily == "Gamma"){
    betaMLE <- coefficients(fit)
    shapeMLE <- as.numeric(gamma.shape(fit)[1])
    rateMLE <- cbind(1, newdata) %*% betaMLE * shapeMLE
    HDCI <- do.call(rbind, lapply(1:nrow(newdata), function(j){
      hdi(qgamma, 1 - alpha, shape = shapeMLE, rate = rateMLE[j, 1])
    })))
  }
  if(confamily == "gaussian"){
    fit = lm(y ~ x1 + I(x1^2) + I(x1^3), data = data)
    betaMLE <- coefficients(fit)
    sdMLE <- summary(fit)$sigma
    meanMLE <- as.numeric(cbind(1, x, x^2, x^3) %*% betaMLE)
  }
}

```

```

HDCI <- do.call(rbind, lapply(1:nrow(newdata), function(j){
  hdi(qnorm, 1 - alpha, sd = sdMLE, mean = meanMLE[j])
}))
}
}

## local coverage prediction regions
output.parametric <- output.nonparametric <-
  output.LS <- output.LSLW <- output.HD <- rep(NA, bins + 1)
if(parametric){
  marginal.parametric <- local.coverage(region = paraCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.parametric <- local.coverage(region = paraCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.parametric <- local.coverage(region = paraCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.parametric <- c(marginal.parametric, local.parametric,
    local.inx.parametric,
    mean(apply(paraCI, 1, diff)),
    absolute.error(y = y, region = paraCI))
}
if(nonparametric){
  marginal.nonparametric <- local.coverage(region = nonparaCI,
    nonparametric = "TRUE", data = data, d = p, bins = 1,
    at.data = "TRUE")
  local.nonparametric <- local.coverage(region = nonparaCI,
    nonparametric = "TRUE", data = data, d = p, bins = bins,
    at.data = "TRUE")
  local.inx.nonparametric <- local.coverage(region = nonparaCI,
    nonparametric = "TRUE", data = data, d = p, bins = 25,
    at.data = "TRUE")
  output.nonparametric <-
    c(marginal.nonparametric, local.nonparametric,
      local.inx.nonparametric,
      area.nonparametric(nonparaCI),
      absolute.error.nonparametric(data = data,
        region = nonparaCI))
}
if(LS){
  marginal.LS <- local.coverage(region = LSCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.LS <- local.coverage(region = LSCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.LS <- local.coverage(region = LSCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.LS <- c(marginal.LS, local.LS, local.inx.LS,
    mean(apply(LSCI, 1, diff)),

```

```

      absolute.error(y = y, region = LSCI))
}
if(LSLW){
  marginal.LSLW <- local.coverage(region = LSLWCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.LSLW <- local.coverage(region = LSLWCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.LSLW <- local.coverage(region = LSLWCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.LSLW <- c(marginal.LSLW, local.LSLW, local.inx.LSLW,
    mean(apply(LSLWCI, 1, diff)),
    absolute.error(y = y, region = LSLWCI))
}
if(HD){
  marginal.HD <- local.coverage(region = HDCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.HD <- local.coverage(region = HDCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.HD <- local.coverage(region = HDCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.HD <- c(marginal.HD, local.HD, local.inx.HD,
    mean(apply(HDCI, 1, diff)),
    absolute.error(y = y, region = HDCI))
}

output <- list(output.parametric = output.parametric,
  output.nonparametric = output.nonparametric,
  output.LS = output.LS,
  output.LSLW = output.LSLW,
  output.HD = output.HD)
output
}

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 5.

```

set.seed(13)
beta <- c(0.5, 1)
n <- 150
bins <- 2
B <- 250
system.time(out.misspec.150.2.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 5))
  })))

```

```
##      user      system elapsed
## 6265.158      57.249 3749.289
```

```
misspec.150.2.5 <- cbind(
  rowMeans(out.misspec.150.2.5, na.rm = TRUE),
  apply(out.misspec.150.2.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 10.

```
system.time(out.misspec.150.2.10 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 10))
  })))
```

```
##      user      system elapsed
## 6352.030      58.194 3757.401
```

```
misspec.150.2.10 <- cbind(
  rowMeans(out.misspec.150.2.10, na.rm = TRUE),
  apply(out.misspec.150.2.10, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 20.

```
system.time(out.misspec.150.2.20 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 20))
  })))
```

```
##      user      system elapsed
## 6415.936      58.042 3757.709
```

```
misspec.150.2.20 <- cbind(
  rowMeans(out.misspec.150.2.20, na.rm = TRUE),
  apply(out.misspec.150.2.20, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 30.

```
system.time(out.misspec.150.2.30 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 30))
  })))

##      user      system elapsed
## 6650.850    58.051 3783.123
```

```
misspec.150.2.30 <- cbind(
  rowMeans(out.misspec.150.2.30, na.rm = TRUE),
  apply(out.misspec.150.2.30, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 40.

```
system.time(out.misspec.150.2.40 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 40))
  })))

##      user      system elapsed
## 6684.993    59.446 3776.350
```

```
misspec.150.2.40 <- cbind(
  rowMeans(out.misspec.150.2.40, na.rm = TRUE),
  apply(out.misspec.150.2.40, 1,
    FUN = function(x){
```

```

sds <- sd(x, na.rm = TRUE)
lengths <- length(which(!is.na(x)))
sds / sqrt(lengths)
}))

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 50.

```

system.time(out.misspec.150.2.50 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 50))
  })))

##      user      system elapsed
## 6695.129    58.823  3771.677

```

```

misspec.150.2.50 <- cbind(
  rowMeans(out.misspec.150.2.50, na.rm = TRUE),
  apply(out.misspec.150.2.50, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 60.

```

system.time(out.misspec.150.2.60 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 60))
  })))

##      user      system elapsed
## 7177.232    61.672  3960.066

```

```

misspec.150.2.60 <- cbind(
  rowMeans(out.misspec.150.2.60, na.rm = TRUE),
  apply(out.misspec.150.2.60, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 70.

```
system.time(out.misspec.150.2.70 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 70))
  })))

##      user      system elapsed
## 7025.012    60.697 3852.552
```

```
misspec.150.2.70 <- cbind(
  rowMeans(out.misspec.150.2.70, na.rm = TRUE),
  apply(out.misspec.150.2.70, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 80.

```
system.time(out.misspec.150.2.80 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 80))
  })))

##      user      system elapsed
## 7108.047    60.454 3859.577
```

```
misspec.150.2.80 <- cbind(
  rowMeans(out.misspec.150.2.80, na.rm = TRUE),
  apply(out.misspec.150.2.80, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 90.


```

system.time(out.misspec.150.2.90 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 90))
  })))

##      user      system elapsed
## 7037.349    60.901 3840.050

```

```

misspec.150.2.90 <- cbind(
  rowMeans(out.misspec.150.2.90, na.rm = TRUE),
  apply(out.misspec.150.2.90, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 250$ iterations when $n = 150$ and shape = 100.

```

system.time(out.misspec.150.2.100 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 100))
  })))

##      user      system elapsed
## 7274.429    60.293 3887.654

```

```

misspec.150.2.100 <- cbind(
  rowMeans(out.misspec.150.2.100, na.rm = TRUE),
  apply(out.misspec.150.2.100, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

Reorganize the output.

```

para.area.misspec.150 <- nonpara.area.misspec.150 <-
  LS.area.misspec.150 <- LSLW.area.misspec.150 <-
  HD.area.misspec.150 <- NULL
para.error.misspec.150 <- nonpara.error.misspec.150 <-
  LS.error.misspec.150 <- LSLW.error.misspec.150 <-

```

```

HD.error.misspec.150 <- NULL
para.marginal.misspec.150 <- nonpara.marginal.misspec.150 <-
  LS.marginal.misspec.150 <- LSLW.marginal.misspec.150 <-
  HD.marginal.misspec.150 <- NULL
para.local.misspec.150 <- nonpara.local.misspec.150 <-
  LS.local.misspec.150 <- LSLW.local.misspec.150 <-
  HD.local.misspec.150 <- NULL
para.inx.misspec.150 <- nonpara.inx.misspec.150 <-
  LS.inx.misspec.150 <- LSLW.inx.misspec.150 <-
  HD.inx.misspec.150 <- NULL
shapes <- c(5, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100)
for(j in shapes ){
  internal.output <- eval(parse(text=paste("misspec.150.2", j, sep = ".")))
  k <- which(shapes == j)
  para.marginal.misspec.150[k] <- as.numeric(internal.output[1, 1])
  para.local.misspec.150[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[2:3, 1])
  para.inx.misspec.150[c((25*(k-1)+1):(25*k))] <- as.numeric(internal.output[4:28,
  para.error.misspec.150[k] <- as.numeric(internal.output[29, 1])
  para.area.misspec.150[k] <- as.numeric(internal.output[30, 1])
  nonpara.marginal.misspec.150[k] <- as.numeric(internal.output[31, 1])
  nonpara.local.misspec.150[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[32:33, 1])
  nonpara.inx.misspec.150[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[34:58, 1])
  nonpara.error.misspec.150[k] <- as.numeric(internal.output[59, 1])
  nonpara.area.misspec.150[k] <- as.numeric(internal.output[60, 1])
  LS.marginal.misspec.150[k] <- as.numeric(internal.output[61, 1])
  LS.local.misspec.150[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[62:63, 1])
  LS.inx.misspec.150[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[64:88, 1])
  LS.error.misspec.150[k] <- as.numeric(internal.output[89, 1])
  LS.area.misspec.150[k] <- as.numeric(internal.output[90, 1])
  LSLW.marginal.misspec.150[k] <- as.numeric(internal.output[91, 1])
  LSLW.local.misspec.150[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[92:93, 1])
  LSLW.inx.misspec.150[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[94:118, 1])
  LSLW.error.misspec.150[k] <- as.numeric(internal.output[119, 1])
  LSLW.area.misspec.150[k] <- as.numeric(internal.output[120, 1])
  HD.marginal.misspec.150[k] <- as.numeric(internal.output[121, 1])
  HD.local.misspec.150[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[122:123, 1])
  HD.inx.misspec.150[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[124:148, 1])
  HD.error.misspec.150[k] <- as.numeric(internal.output[149, 1])

```

```
HD.area.misspec.150[k] <- as.numeric(internal.output[150, 1])
}
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and $\text{shape} = 5$.

```
set.seed(13)
n <- 250
bins <- 3
system.time(out.misspec.250.3.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 5))
  })))

##      user      system elapsed
## 2392.483    14.124 1461.153
```

```
misspec.250.3.5 <- cbind(
  rowMeans(out.misspec.250.3.5, na.rm = TRUE),
  apply(out.misspec.250.3.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and $\text{shape} = 20$.

```
system.time(out.misspec.250.3.20 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 20))
  })))

##      user      system elapsed
## 2296.508    14.109 1429.061
```

```
misspec.250.3.20 <- cbind(
  rowMeans(out.misspec.250.3.20, na.rm = TRUE),
  apply(out.misspec.250.3.20, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 40.

```
system.time(out.misspec.250.3.40 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 40))
  })))

##      user      system elapsed
## 2282.950    13.735 1413.949
```

```
misspec.250.3.40 <- cbind(
  rowMeans(out.misspec.250.3.40, na.rm = TRUE),
  apply(out.misspec.250.3.40, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 60.

```
system.time(out.misspec.250.3.60 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 60))
  })))

##      user      system elapsed
## 2414.106    13.831 1435.658
```

```
misspec.250.3.60 <- cbind(
  rowMeans(out.misspec.250.3.60, na.rm = TRUE),
  apply(out.misspec.250.3.60, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })))
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 80.

```

system.time(out.misspec.250.3.80 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 80))
  })))

##      user      system elapsed
## 2395.971    13.535 1425.224

```

```

misspec.250.3.80 <- cbind(
  rowMeans(out.misspec.250.3.80, na.rm = TRUE),
  apply(out.misspec.250.3.80, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 250$ and shape = 100.

```

system.time(out.misspec.250.3.100 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 100))
  })))

##      user      system elapsed
## 2391.049    12.727 1422.181

```

```

misspec.250.3.100 <- cbind(
  rowMeans(out.misspec.250.3.100, na.rm = TRUE),
  apply(out.misspec.250.3.100, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

Reorganize the output.

```

para.area.misspec.250 <- nonpara.area.misspec.250 <-
  LS.area.misspec.250 <- LSLW.area.misspec.250 <-
  HD.area.misspec.250 <- NULL
para.error.misspec.250 <- nonpara.error.misspec.250 <-
  LS.error.misspec.250 <- LSLW.error.misspec.250 <-

```

```

HD.error.misspec.250 <- NULL
para.marginal.misspec.250 <- nonpara.marginal.misspec.250 <-
  LS.marginal.misspec.250 <- LSLW.marginal.misspec.250 <-
  HD.marginal.misspec.250 <- NULL
para.local.misspec.250 <- nonpara.local.misspec.250 <-
  LS.local.misspec.250 <- LSLW.local.misspec.250 <-
  HD.local.misspec.250 <- NULL
para.inx.misspec.250 <- nonpara.inx.misspec.250 <-
  LS.inx.misspec.250 <- LSLW.inx.misspec.250 <-
  HD.inx.misspec.250 <- NULL
shapes <- c(5, 20, 40, 60, 80, 100)
for(j in shapes ){
  internal.output <- eval(parse(text=paste("misspec.250.3", j, sep = ".")))
  k <- which(shapes == j)
  para.marginal.misspec.250[k] <- as.numeric(internal.output[1, 1])
  para.local.misspec.250[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[2:4, 1])
  para.inx.misspec.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[5:29, 1])
  para.error.misspec.250[k] <- as.numeric(internal.output[30, 1])
  para.area.misspec.250[k] <- as.numeric(internal.output[31, 1])
  nonpara.marginal.misspec.250[k] <- as.numeric(internal.output[32, 1])
  nonpara.local.misspec.250[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[33:35, 1])
  nonpara.inx.misspec.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[36:60, 1])
  nonpara.error.misspec.250[k] <- as.numeric(internal.output[61, 1])
  nonpara.area.misspec.250[k] <- as.numeric(internal.output[62, 1])
  LS.marginal.misspec.250[k] <- as.numeric(internal.output[63, 1])
  LS.local.misspec.250[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[64:66, 1])
  LS.inx.misspec.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[67:91, 1])
  LS.error.misspec.250[k] <- as.numeric(internal.output[92, 1])
  LS.area.misspec.250[k] <- as.numeric(internal.output[93, 1])
  LSLW.marginal.misspec.250[k] <- as.numeric(internal.output[94, 1])
  LSLW.local.misspec.250[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[95:97, 1])
  LSLW.inx.misspec.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[98:122, 1])
  LSLW.error.misspec.250[k] <- as.numeric(internal.output[123, 1])
  LSLW.area.misspec.250[k] <- as.numeric(internal.output[124, 1])
  HD.marginal.misspec.250[k] <- as.numeric(internal.output[125, 1])
  HD.local.misspec.250[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[126:128, 1])
  HD.inx.misspec.250[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[129:153, 1])
}

```

```

HD.error.misspec.250[k] <- as.numeric(internal.output[154, 1])
HD.area.misspec.250[k] <- as.numeric(internal.output[155, 1])
}

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and $\text{shape} = 5$.

```

set.seed(13)
n <- 500
bins <- 3
system.time(out.misspec.500.3.5 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 5))
  })))

##      user      system elapsed
## 6925.010    19.527 4348.905

```

```

misspec.500.3.5 <- cbind(
  rowMeans(out.misspec.500.3.5, na.rm = TRUE),
  apply(out.misspec.500.3.5, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and $\text{shape} = 20$.

```

system.time(out.misspec.500.3.20 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 20))
  })))

##      user      system elapsed
## 6726.821    19.619 4257.396

```

```

misspec.500.3.20 <- cbind(
  rowMeans(out.misspec.500.3.20, na.rm = TRUE),
  apply(out.misspec.500.3.20, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 40.

```
system.time(out.misspec.500.3.40 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 40))
  })))

##      user      system elapsed
## 6859.403    19.266  4252.117
```

```
misspec.500.3.40 <- cbind(
  rowMeans(out.misspec.500.3.40, na.rm = TRUE),
  apply(out.misspec.500.3.40, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 60.

```
system.time(out.misspec.500.3.60 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 60))
  })))

##      user      system elapsed
## 7253.875    19.272  4346.840
```

```
misspec.500.3.60 <- cbind(
  rowMeans(out.misspec.500.3.60, na.rm = TRUE),
  apply(out.misspec.500.3.60, 1,
  FUN = function(x){
    sds <- sd(x, na.rm = TRUE)
    lengths <- length(which(!is.na(x)))
    sds / sqrt(lengths)
  })
)
```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 80.


```

system.time(out.misspec.500.3.80 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 80))
  })))

##      user      system elapsed
## 7358.336    19.303  4335.417

```

```

misspec.500.3.80 <- cbind(
  rowMeans(out.misspec.500.3.80, na.rm = TRUE),
  apply(out.misspec.500.3.80, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$ and shape = 100.

```

system.time(out.misspec.500.3.100 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(misspec.simulator(beta = beta, n = n,
      bins = bins, shape = 100))
  })))

##      user      system elapsed
## 7433.921    19.266  4345.724

```

```

misspec.500.3.100 <- cbind(
  rowMeans(out.misspec.500.3.100, na.rm = TRUE),
  apply(out.misspec.500.3.100, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

Reorganize the output.

```

para.area.misspec.500 <- nonpara.area.misspec.500 <-
  LS.area.misspec.500 <- LSLW.area.misspec.500 <-
  HD.area.misspec.500 <- NULL
para.error.misspec.500 <- nonpara.error.misspec.500 <-
  LS.error.misspec.500 <- LSLW.error.misspec.500 <-

```

```

HD.error.misspec.500 <- NULL
para.marginal.misspec.500 <- nonpara.marginal.misspec.500 <-
  LS.marginal.misspec.500 <- LSLW.marginal.misspec.500 <-
  HD.marginal.misspec.500 <- NULL
para.local.misspec.500 <- nonpara.local.misspec.500 <-
  LS.local.misspec.500 <- LSLW.local.misspec.500 <-
  HD.local.misspec.500 <- NULL
para.inx.misspec.500 <- nonpara.inx.misspec.500 <-
  LS.inx.misspec.500 <- LSLW.inx.misspec.500 <-
  HD.inx.misspec.500 <- NULL
shapes <- c(5, 20, 40, 60, 80, 100)
for(j in shapes ){
  internal.output <- eval(parse(text=paste("misspec.500.3", j, sep = ".")))
  k <- which(shapes == j)
  para.marginal.misspec.500[k] <- as.numeric(internal.output[1, 1])
  para.local.misspec.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[2:4, 1])
  para.inx.misspec.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[5:29, 1])
  para.error.misspec.500[k] <- as.numeric(internal.output[30, 1])
  para.area.misspec.500[k] <- as.numeric(internal.output[31, 1])
  nonpara.marginal.misspec.500[k] <- as.numeric(internal.output[32, 1])
  nonpara.local.misspec.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[33:35, 1])
  nonpara.inx.misspec.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[36:60, 1])
  nonpara.error.misspec.500[k] <- as.numeric(internal.output[61, 1])
  nonpara.area.misspec.500[k] <- as.numeric(internal.output[62, 1])
  LS.marginal.misspec.500[k] <- as.numeric(internal.output[63, 1])
  LS.local.misspec.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[64:66, 1])
  LS.inx.misspec.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[67:91, 1])
  LS.error.misspec.500[k] <- as.numeric(internal.output[92, 1])
  LS.area.misspec.500[k] <- as.numeric(internal.output[93, 1])
  LSLW.marginal.misspec.500[k] <- as.numeric(internal.output[94, 1])
  LSLW.local.misspec.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[95:97, 1])
  LSLW.inx.misspec.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[98:122, 1])
  LSLW.error.misspec.500[k] <- as.numeric(internal.output[123, 1])
  LSLW.area.misspec.500[k] <- as.numeric(internal.output[124, 1])
  HD.marginal.misspec.500[k] <- as.numeric(internal.output[125, 1])
  HD.local.misspec.500[c((bins*(k-1)+1):(bins*k))] <-
    as.numeric(internal.output[126:128, 1])
  HD.inx.misspec.500[c((25*(k-1)+1):(25*k))] <-
    as.numeric(internal.output[129:153, 1])
}

```

```
HD.error.misspec.500[k] <- as.numeric(internal.output[154, 1])  
HD.area.misspec.500[k] <- as.numeric(internal.output[155, 1])  
}
```

5.2 Results

Results from our simulations are depicted in Figures 9-14. For all five prediction regions we depict the estimated area, prediction error, and local coverage probabilities with respect to binning in Figures 9, 11, and 13 for $n = 150, 250$, and 500 respectively. For all five prediction regions we depict the local coverage probabilities across x in Figures 10, 12, and 14 for $n = 150, 250$, and 500 respectively.

In these simulations, we expect for the LSLW conformal prediction region to perform well. The model misspecification is modest, the Gamma data appears to be almost symmetric, albeit heterogenous, about a cubic mean function. From these simulations we see that the parametric conformal prediction region is similar to the LSLW prediction region in area and prediction error. Both the parametric and LSLW conformal prediction regions possess finite-sample, albeit slightly conservative, local validity with respect to binning and they both give almost nominal finite-sample conditional validity across the support. Moreover, these prediction regions visually fit the data well as seen in Section 7.2. The nonparametric conformal prediction region gives closer to nominal coverage than both the parametric and LSLW conformal prediction regions, and it possesses finite-sample local validity with respect to binning. However, this prediction region is larger, gives higher prediction error than the parametric and LSLW conformal prediction regions, and does not visually fit the data well for most simulation settings as seen in Section 7.2. The misspecified HD and LS conformal prediction regions provide extreme undercoverage at small values of x and extreme overcoverage at large values of x . These prediction regions are also larger, they give higher prediction error than the parametric and LSLW conformal prediction regions, and the LS conformal prediction region does not visually fit the data well for most simulation settings as seen in Section 7.2.




Figure 9: This figure compares the performance of the misspecified parametric, nonparametric, least squares, and least squares locally weighted conformal prediction region and the misspecified highest density prediction region when $n = 150$ and the number of bins equals 2. The specific diagnostics used to compare these prediction regions is the area (top-left panel), prediction error (top-right panel), and the coverage probability with respect to binning (bottom row) across shape parameter values. The average of 250 Monte Carlo samples at each shape parameter value in these simulation settings form the lines that are depicted in this figure.

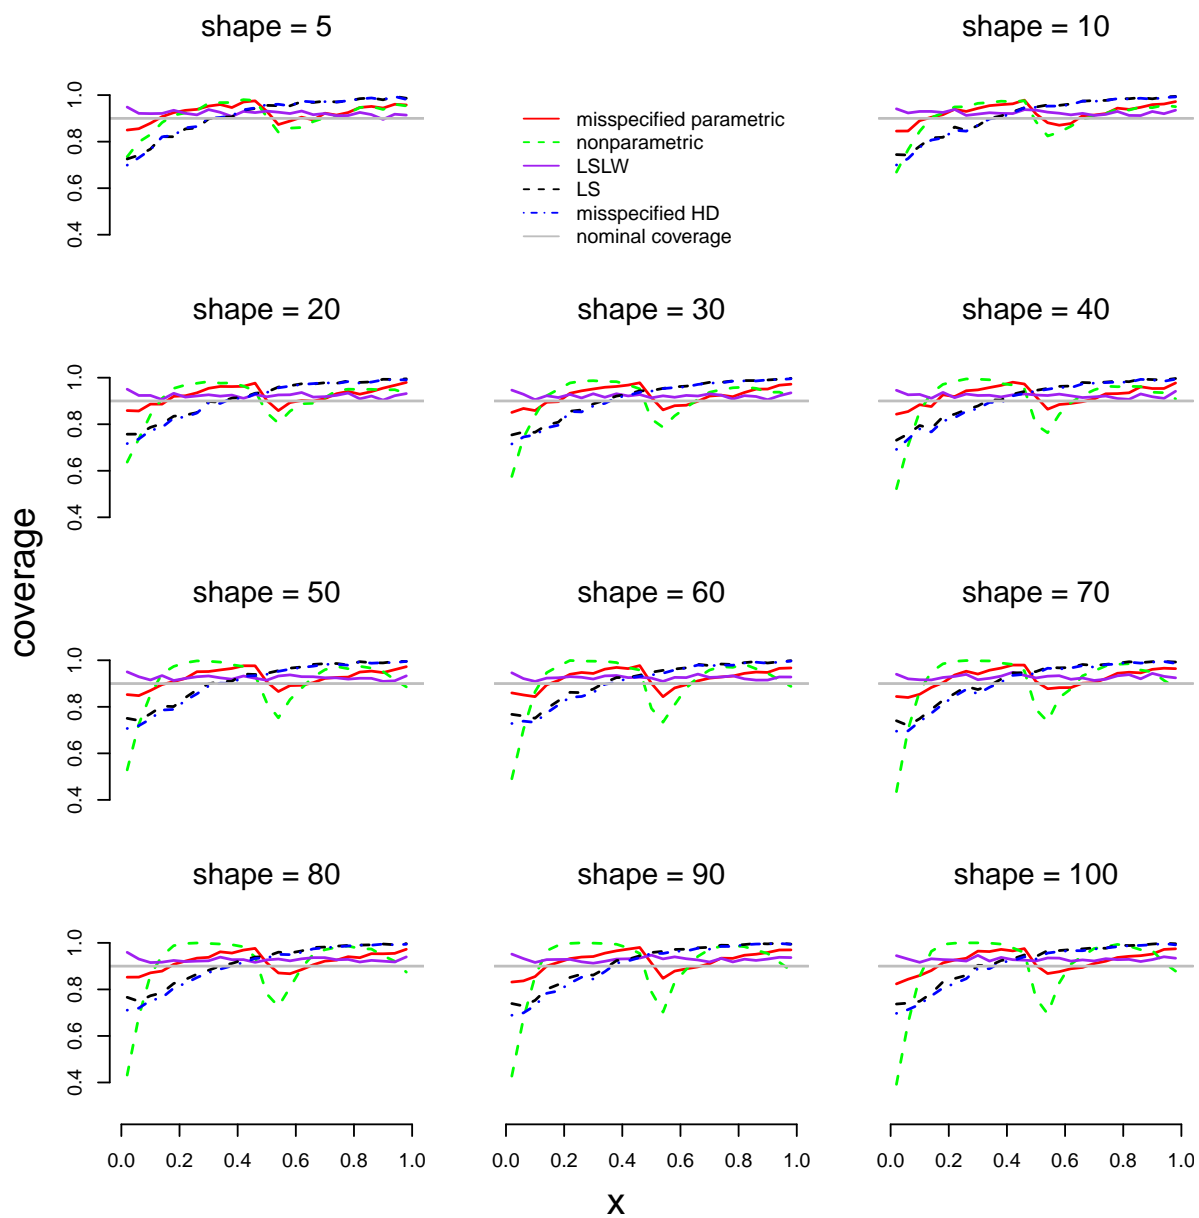


Figure 10: Plot of the estimated coverage probabilities of prediction regions across x and shape parameter values when the model is misspecified, $n = 150$, and the number of bins is equal to 2.




Figure 11: This figure compares the performance of the misspecified parametric, nonparametric, least squares, and least squares locally weighted conformal prediction region and the misspecified highest density prediction region when $n = 250$ and the number of bins equals 3. The specific diagnostics used to compare these prediction regions is the area (top-left panel), prediction error (top-right panel), and the coverage probability with respect to binning (bottom row) across shape parameter values. The average of 50 Monte Carlo samples at each shape parameter value in these simulation settings form the lines that are depicted in this figure.




Figure 12: Plot of the estimated coverage probabilities of prediction regions across x and shape parameter values when the model is misspecified, $n = 250$, and the number of bins is equal to 3.

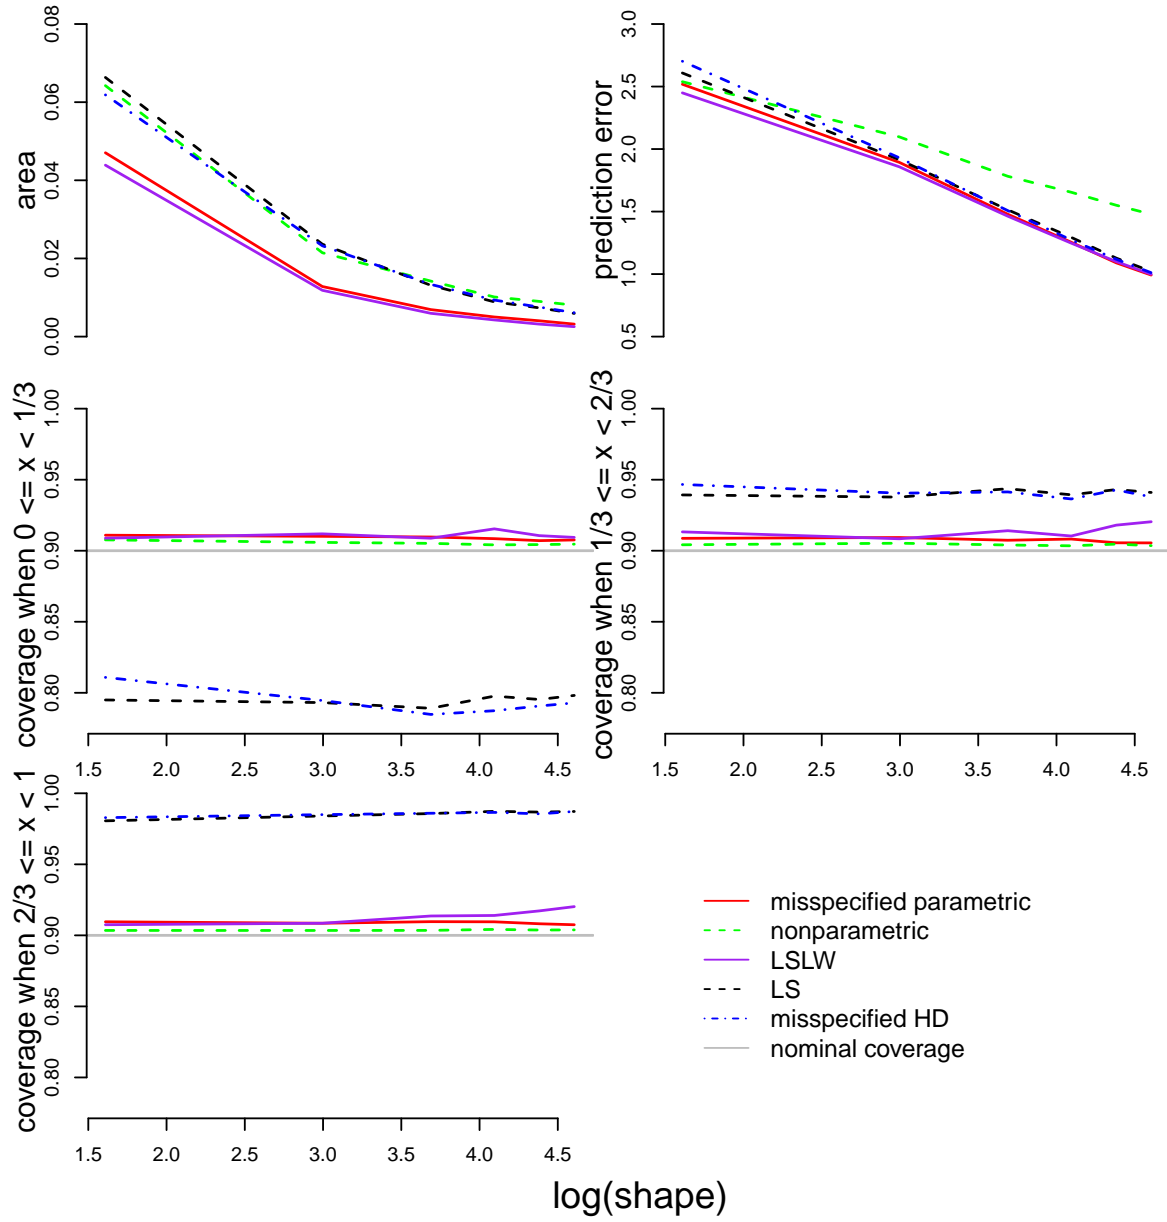


Figure 13: This figure compares the performance of the misspecified parametric, nonparametric, least squares, and least squares locally weighted conformal prediction region and the misspecified highest density prediction region when $n = 500$ and the number of bins equals 3. The specific diagnostics used to compare these prediction regions is the area (top-left panel), prediction error (top-right panel), and the coverage probability with respect to binning (bottom row) across shape parameter values. The average of 50 Monte Carlo samples at each shape parameter value in these simulation settings form the lines that are depicted in this figure.




Figure 14: Plot of the estimated coverage probabilities of prediction regions across x and shape parameter values when the model is misspecified, $n = 500$, and the number of bins is equal to 3.

6 Linear Regression Simulations

In this Section, we compare the parametric conformal prediction region, the nonparametric conformal prediction region, the LSLW conformal prediction region, the LS conformal prediction region, and the HD prediction region for linear regression models with normal errors and constant variance. We specify that $\beta = (2, 5)^T$ and the standard deviation of the errors about the mean function as $\sigma = 1$. We consider sample sizes of $n \in \{150, 250, 500\}$. When $n = 150$ we build the parametric and nonparametric conformal prediction regions using 2 bins. When $n = 250, 500$ we build the parametric and nonparametric conformal prediction regions using 3 bins. These number of bin choices correspond to the bin width asymptotics of Lei and Wasserman [2014].

6.1 Simulations

The following function computes our diagnostic measures for the five prediction regions under investigation in the univariate case where data is generated from a Gaussian regression model, i.e. $Y \sim N(X'\beta, \sigma^2)$ and $X \sim U(0, 1)$.

```
regression.simulator <- function(n = 500, alpha = 0.10, beta,
  bins = 3, sd = 1, parametric = TRUE, nonparametric = TRUE,
  LS = TRUE, LSLW = TRUE, HD = TRUE, cores = 6){

  p <- d <- length(beta) - 1
  x <- matrix(runif(n), ncol = p)
  y <- rep(0, n)
  data <- NULL

  ## set up partition
  if(class(bins) == "NULL"){
    wn <- min(1/ floor(1 / (log(n)/n)^(1/(d+3))), 1/2)
    bins <- 1 / wn
  }

  ## generate the data
  mu <- cbind(1, x) %*% beta
  y <- rnorm(n = n, mean = mu, sd = sd)
  data <- data.frame(y = y, x = x)
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")

  ## fit the misspecified cubic regression model
  fit <- glm(y ~ x1, family = "gaussian", data = data)
  paraCI <- nonparaCI <- LSCI <- LSLWCI <- HDCI <- NULL
  formula <- fit$formula
  newdata <- data
  respname <- all.vars(formula)[1]
  newdata <- newdata[, !(colnames(data) %in% respname)]
  newdata <- as.matrix(newdata)
```

```

## obtain the prediction regions
if(parametric){
  cpred <- conformal.glm(fit, parametric = TRUE,
    nonparametric = FALSE, alpha = alpha,
    bins = bins, cores = cores)
  paraCI <- cpred$paraconformal
}
if(nonparametric){
  cpred <- conformal.glm(fit, parametric = FALSE,
    nonparametric = TRUE, alpha = alpha,
    bins = bins, cores = cores)
  nonparaCI <- cpred$nonparaconformal
}
if(LS){
  p1.tibs <- conformal.pred(x = x, y = y, x0 = x,
    train.fun = train.fun, predict.fun = predict.fun,
    alpha = alpha)
  LSCI <- cbind(p1.tibs$lo, p1.tibs$up)
}
if(LSLW){
  regression.model <- lm(y ~ x)
  abs.resid <- abs(regression.model$resid)
  smooth.call <- smooth.spline(x, abs.resid,
    nknots = 10)
  lambda <- smooth.call$lambda
  df <- smooth.call$df
  mad.train.fun <- function(x, y, out = NULL){
    smooth.spline(x[, 1], y, lambda = lambda,
      df = df, nknots = 10)
  }
  p2.tibs <- conformal.pred(x = x, y = y, x0 = x,
    train.fun = train.fun, predict.fun = predict.fun,
    mad.train.fun = mad.train.fun,
    mad.predict.fun = mad.predict.fun,
    alpha = alpha)
  LSLWCI <- cbind(p2.tibs$lo, p2.tibs$up)
}
if(HD){
  fit = lm(y ~ x, data = data)
  betaMLE <- coefficients(fit)
  sdMLE <- summary(fit)$sigma
  meanMLE <- as.numeric(cbind(1, x) %*% betaMLE)
  HDCI <- do.call(rbind, lapply(1:nrow(newdata), function(j){
    hdi(qnorm, 1 - alpha, sd = sdMLE, mean = meanMLE[j])
  })))
}

```

```

## local coverage prediction regions
output.parametric <- output.nonparametric <-
  output.LS <- output.LSLW <- output.HD <- rep(NA, bins + 1)
if(parametric){
  marginal.parametric <- local.coverage(region = paraCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.parametric <- local.coverage(region = paraCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.parametric <- local.coverage(region = paraCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.parametric <- c(marginal.parametric, local.parametric,
    local.inx.parametric,
    mean(apply(paraCI, 1, diff)),
    absolute.error(y = y, region = paraCI))
}
if(nonparametric){
  marginal.nonparametric <- local.coverage(region = nonparaCI,
    nonparametric = "TRUE", data = data, d = p, bins = 1,
    at.data = "TRUE")
  local.nonparametric <- local.coverage(region = nonparaCI,
    nonparametric = "TRUE", data = data, d = p, bins = bins,
    at.data = "TRUE")
  local.inx.nonparametric <- local.coverage(region = nonparaCI,
    nonparametric = "TRUE", data = data, d = p, bins = 25,
    at.data = "TRUE")
  output.nonparametric <-
    c(marginal.nonparametric, local.nonparametric,
      local.inx.nonparametric,
      area.nonparametric(nonparaCI),
      absolute.error.nonparametric(data = data,
        region = nonparaCI))
}
if(LS){
  marginal.LS <- local.coverage(region = LSCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.LS <- local.coverage(region = LSCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.LS <- local.coverage(region = LSCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.LS <- c(marginal.LS, local.LS, local.inx.LS,
    mean(apply(LSCI, 1, diff)),
    absolute.error(y = y, region = LSCI))
}
if(LSLW){
  marginal.LSLW <- local.coverage(region = LSLWCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.LSLW <- local.coverage(region = LSLWCI,

```

```

    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.LSLW <- local.coverage(region = LSLWCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.LSLW <- c(marginal.LSLW, local.LSLW, local.inx.LSLW,
    mean(apply(LSLWCI, 1, diff)),
    absolute.error(y = y, region = LSLWCI))
}
if(HD){
  marginal.HD <- local.coverage(region = HDCI,
    data = data, d = p, bins = 1, at.data = "TRUE")
  local.HD <- local.coverage(region = HDCI,
    data = data, d = p, bins = bins, at.data = "TRUE")
  local.inx.HD <- local.coverage(region = HDCI,
    data = data, d = p, bins = 25, at.data = "TRUE")
  output.HD <- c(marginal.HD, local.HD, local.inx.HD,
    mean(apply(HDCI, 1, diff)),
    absolute.error(y = y, region = HDCI))
}

output <- list(output.parametric = output.parametric,
  output.nonparametric = output.nonparametric,
  output.LS = output.LS,
  output.LSLW = output.LSLW,
  output.HD = output.HD)
output
}

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 150$.

```

set.seed(13)
beta <- c(2, 5)
n <- 150
bins <- 2
B <- 50
system.time(out.regression.150.2 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(regression.simulator(beta = beta, n = n,
      bins = bins))
  })))

##      user      system elapsed
## 1509.298      13.117    901.467

```

```

regression.150.2 <- cbind(
  rowMeans(out.regression.150.2, na.rm = TRUE),
  apply(out.regression.150.2, 1,

```

```

FUN = function(x){
  sds <- sd(x, na.rm = TRUE)
  lengths <- length(which(!is.na(x)))
  sds / sqrt(lengths)
})

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 150$.

```

n <- 250
bins <- 3
system.time(out.regression.250.3 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(regression.simulator(beta = beta, n = n,
      bins = bins))
  })))

##      user      system elapsed
## 2945.050      18.113 1846.314

```

```

regression.250.3 <- cbind(
  rowMeans(out.regression.250.3, na.rm = TRUE),
  apply(out.regression.250.3, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))
      sds / sqrt(lengths)
    })
)

```

The following performs our Monte Carlo simulation of $B = 50$ iterations when $n = 500$.

```

n <- 500
system.time(out.regression.500.3 <- do.call(cbind,
  lapply(1:B, FUN = function(j){
    unlist(regression.simulator(beta = beta, n = n,
      bins = bins))
  })))

##      user      system elapsed
## 7285.245      19.752 4517.354

```

```

regression.500.3 <- cbind(
  rowMeans(out.regression.500.3, na.rm = TRUE),
  apply(out.regression.500.3, 1,
    FUN = function(x){
      sds <- sd(x, na.rm = TRUE)
      lengths <- length(which(!is.na(x)))

```

```
sds / sqrt(lengths)
}))
```

6.2 Results

Results from our simulations are depicted in Table 2 and Figure 15. This table and figure depicts the estimated area, prediction error, and local coverage probabilities for all five considered prediction regions.

In these simulations, errors about the estimated mean function are symmetric and homogeneous across the support. We therefore expect for the LS and LSLW conformal prediction regions to perform nearly as well as the oracle HD prediction region. These prediction regions also exhibit finite-sample marginal validity, local validity with respect to binning, and near conditional validity across the support. The parametric conformal prediction region is similar to the LS and LSLW conformal prediction regions and the HD prediction region in area, prediction error, finite-sample coverage properties, and appearance. However, the parametric conformal prediction region is slightly larger and gives more conservative coverage than these other prediction regions. The nonparametric conformal prediction region is larger and gives larger prediction errors than the other prediction regions. It also appears to not visually fit the data well while the others do as seen in Section 7.3.

		parametric conformal	nonparametric conformal	LS conformal	LSLW conformal	HD region
$n = 150$	marginal coverage	0.922 (0.0007)	0.916 (0.001)	0.913 (0.001)	0.918 (0.0011)	0.904 (0.0023)
	local coverage when $0 < x < 1/2$	0.922 (0.0011)	0.916 (0.0014)	0.912 (0.0035)	0.921 (0.0029)	0.903 (0.0042)
	local coverage when $1/2 \leq x < 1$	0.922 (0.001)	0.916 (0.0016)	0.914 (0.0036)	0.915 (0.003)	0.904 (0.0036)
	area	3.521 (0.0366)	4.258 (0.0466)	3.361 (0.0291)	3.385 (0.0322)	3.28 (0.0263)
	prediction error	0.021 (0.0011)	0.033 (0.0023)	0.027 (0.0015)	0.023 (0.0013)	0.03 (0.0013)
$n = 250$	marginal coverage	0.918 (0.0006)	0.915 (0.0007)	0.908 (0.0006)	0.911 (0.0007)	0.901 (0.0014)
	local coverage when $0 < x < 1/3$	0.919 (0.001)	0.915 (0.0014)	0.91 (0.0036)	0.912 (0.0034)	0.905 (0.0039)
	local coverage when $1/3 \leq x < 2/3$	0.917 (0.0007)	0.917 (0.0014)	0.912 (0.0035)	0.913 (0.0028)	0.907 (0.0039)
	local coverage when $2/3 \leq x < 1$	0.918 (0.001)	0.915 (0.0015)	0.904 (0.0042)	0.907 (0.0032)	0.893 (0.0046)
	area	3.518 (0.0291)	3.822 (0.0338)	3.36 (0.0233)	3.363 (0.0261)	3.296 (0.0218)
$n = 500$	prediction error	0.021 (0.0013)	0.028 (0.0017)	0.027 (0.0013)	0.025 (0.0013)	0.029 (0.0013)
	marginal coverage	0.908 (0.0002)	0.907 (0.0004)	0.905 (0.0005)	0.906 (0.0005)	0.903 (0.0013)
	local coverage when $0 < x < 1/3$	0.909 (0.0004)	0.907 (0.0008)	0.907 (0.0032)	0.906 (0.0025)	0.905 (0.0034)
	local coverage when $1/3 \leq x < 2/3$	0.908 (0.0004)	0.906 (0.0006)	0.907 (0.0026)	0.906 (0.0021)	0.907 (0.0029)
	local coverage when $2/3 \leq x < 1$	0.909 (0.0005)	0.908 (0.0007)	0.901 (0.0024)	0.905 (0.0021)	0.898 (0.0028)
	area	3.369 (0.0211)	3.715 (0.0187)	3.305 (0.0206)	3.314 (0.0202)	3.287 (0.0169)
	prediction error	0.026 (0.0013)	0.033 (0.0013)	0.029 (0.0014)	0.028 (0.0014)	0.03 (0.0011)

Table 2: Diagnostics for conformal prediction regions for linear regression models with normal errors and constant variance. Local and marginal coverage properties, areas, and prediction errors are presented for the parametric conformal prediction region (third column), nonparametric conformal prediction region (fourth column), LS conformal prediction region (fifth column), LSLW conformal prediction region (sixth column), and HD prediction region (seventh column). Standard errors are in parentheses.




Figure 15: Plot of the estimated coverage probabilities of prediction regions across x and sample sizes.

7 Example plots of prediction regions

In this section we construct prediction regions corresponding to the simulations and results in Sections 4 through 7.3. In Gamma analyses we generate a dataset for each shape parameter considered with $n = 150$ and in regression analyses we generate a dataset for all sample sizes considered. For each of these datasets we depict the parametric, nonparametric, LS, LSLW conformal prediction regions over the observed data to visually assess the appropriateness of each prediction region. The findings from these figures are consistent with the findings from our numerical diagnostics. The parametric conformal prediction region gives visually natural bounds for the observed data when the model is correctly specified in small to moderate sample sizes and is appropriate when the model is misspecified. The LSLW conformal prediction region gives visually natural bounds for the observed data when the model is correctly specified, is appropriate under mild model misspecification, and is ill-fitting in settings where deviations about an estimated mean function are clearly not symmetric. The nonparametric conformal prediction region is coarse and is larger than necessary when the mean function is steep relative to its variability. The LS conformal prediction performs well when deviations about the estimated mean function are symmetric and is sensitive to mild departures from that setting. This prediction region is seen to provide overcoverage (undercoverage) in regions of the predictor space where variability about the estimated mean function is relatively small (large).

7.1 Plots corresponding to Section 4

The following function makes the conformal prediction region plots for our Gamma simulations under correct parametric conformal model specification.

```
makeplots <- function(n = 500, bins = 3, shape = 1, beta){  
  
  ## generate random data  
  set.seed(13)  
  alpha <- 0.10  
  cores <- 6  
  p <- k <- length(beta) - 1  
  x <- matrix(runif(n), ncol = p)  
  rate <- (cbind(1, x) %*% beta) * shape  
  y <- rgamma(n = n, shape = shape, rate = rate)  
  y <- y / sd(y)  
  data <- data.frame(y = y, x = x)  
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")  
  
  ## fit the Gamma regression model  
  fit <- glm(y ~ x1, family = "Gamma", data = data)  
  formula <- fit$formula  
  newdata <- data  
  resname <- all.vars(formula)[1]  
  newdata <- newdata[, !(colnames(data) %in% resname)]  
  newdata <- as.matrix(newdata)  
  
  ## obtain parametric and nonparametric conformal  
  ## prediction regions  
  cpred <- conformal.glm(fit, parametric = TRUE,
```

```

    nonparametric = FALSE, alpha = alpha,
    bins = bins, cores = cores)
paraCI <- cpred$paraconformal

cpred <- conformal.glm(fit, parametric = FALSE,
  nonparametric = TRUE, alpha = alpha,
  bins = bins, cores = cores)
nonparaCI <- cpred$nonparaconformal

## obtain LS conformal prediction region
p1.tibs <- conformal.pred(x = cbind(x, x^2, x^3), y = y,
  x0 = cbind(x, x^2, x^3), train.fun = train.fun,
  predict.fun = predict.fun, alpha = alpha,
  grid.method = "linear")
LSICI <- cbind(p1.tibs$lo, p1.tibs$up)

## obtain LSLW conformal prediction region
cubic.model <- lm(y ~ x + I(x^2) + I(x^3))
abs.resid <- abs(cubic.model$resid)
smooth.call <- smooth.spline(x, abs.resid,
  nknots = 10)
lambda <- smooth.call$lambda
df <- smooth.call$df
mad.train.fun <- function(x, y, out = NULL){
  smooth.spline(x[, 1], y, lambda = lambda,
    df = df, nknots = 10)
}
p2.tibs <- conformal.pred(x = cbind(x, x^2, x^3), y = y,
  x0 = cbind(x, x^2, x^3),
  train.fun = train.fun, predict.fun = predict.fun,
  mad.train.fun = mad.train.fun,
  mad.predict.fun = mad.predict.fun, alpha = alpha)
LSLWCI <- cbind(p2.tibs$lo, p2.tibs$up)

## create plots
ix <- sort(x, index.return = TRUE)$ix
par(mfrow = c(1,4), oma = c(4,4,0,0), mar = c(1,1,1,1))

## parametric conformal
plot.new()
plot.window(xlim = c(0,1), ylim = c(min(y)-1, max(y)+1))
points(x, y, pch = 19, col = rgb(0,0,0,alpha=0.2))
lines(x[ix], paraCI[ix, 1], col = "red", lwd = 1.5)
lines(x[ix], paraCI[ix, 2], col = "red", lwd = 1.5)
axis(1); axis(2)

## LSLW conformal

```

```

plot.new()
plot.window(xlim = c(0,1), ylim = c(min(y)-1, max(y)+1))
points(x, y, pch = 19, col = rgb(0,0,0,alpha=0.2))
lines(x[ix], LSLWCI[ix, 1], col = "red", lwd = 1.5)
lines(x[ix], LSLWCI[ix, 2], col = "red", lwd = 1.5)
axis(1)

## nonparametric conformal
plot.nonparametric(region = nonparaCI, x = x,
  y = y, bins = bins)
axis(1)

## LS conformal
plot.new()
plot.window(xlim = c(0,1), ylim = c(min(y)-1, max(y)+1))
points(x, y, pch = 19, col = rgb(0,0,0,alpha=0.2))
lines(x[ix], LSCI[ix, 1], col = "red", lwd = 1.5)
lines(x[ix], LSCI[ix, 2], col = "red", lwd = 1.5)
axis(1)

mtext("x", side = 1, cex = 1, line = 2, outer = TRUE)
mtext("y", side = 2, cex = 1, line = 2, outer = TRUE)
}

```

```

beta <- c(1.25, -1)
makeplots(n = 150, shape = 0.5, bins = 2, beta = beta)

```




Figure 16: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 0.5$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 0.75, bins = 2, beta = beta)
```




Figure 17: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 0.75$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 1, bins = 2, beta = beta)
```

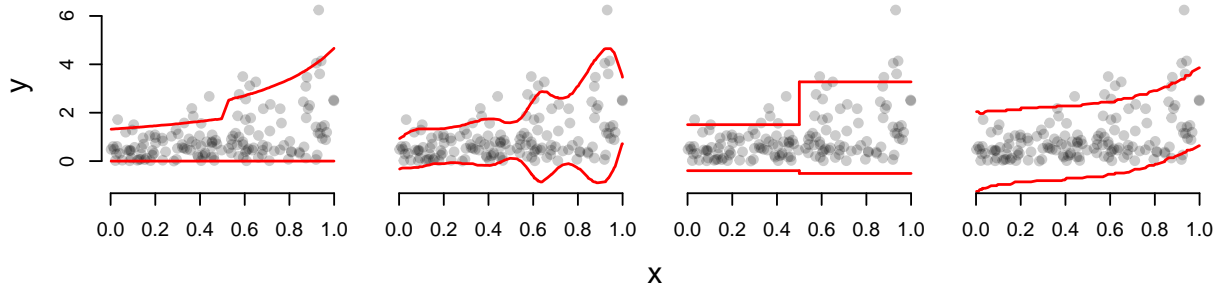


Figure 18: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 1$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 2, bins = 2, beta = beta)
```




Figure 19: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 2$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 5, bins = 2, beta = beta)
```




Figure 20: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 2$, and $\text{bins} = 5$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 7, bins = 2, beta = beta)
```




Figure 21: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 7$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 10, bins = 2, beta = beta)
```




Figure 22: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 10$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 12, bins = 2, beta = beta)
```




Figure 23: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 12$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 15, bins = 2, beta = beta)
```




Figure 24: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 15$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).


```
makeplots(n = 150, shape = 20, bins = 2, beta = beta)
```




Figure 25: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 20$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 25, bins = 2, beta = beta)
```




Figure 26: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 25$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 50, bins = 2, beta = beta)
```




Figure 27: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 50$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots(n = 150, shape = 100, bins = 2, beta = beta)
```




Figure 28: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 100$, and $\text{bins} = 2$. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

7.2 Plots corresponding to Section 5

The following function makes the conformal prediction region plots for our Gamma simulations under incorrect parametric conformal model specification.

```
makeplots.misspec <- function(n = 500, bins = 3, shape = 1, beta){  
  
  ## generate random data  
  set.seed(13)  
  alpha <- 0.10  
  cores <- 6  
  p <- k <- length(beta) - 1  
  x <- matrix(runif(n), ncol = p)  
  rate <- (cbind(1, x) %*% beta) * shape  
  y <- rgamma(n = n, shape = shape, rate = rate)  
  y <- y / sd(y)  
  data <- data.frame(y = y, x = x)  
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")  
  
  ## fit the Gamma regression model  
  fit <- glm(y ~ x1 + I(x1^2) + I(x1^2),  
    family = "gaussian", data = data)  
  formula <- fit$formula  
  newdata <- data  
  respname <- all.vars(formula)[1]  
  newdata <- newdata[, !(colnames(data) %in% respname)]  
  newdata <- as.matrix(newdata)  
  
  ## obtain parametric and nonparametric conformal  
  ## prediction regions  
  cpred <- conformal.glm(fit, parametric = TRUE,  
    nonparametric = FALSE, alpha = alpha,  
    bins = bins, cores = cores)  
  paraCI <- cpred$paraconformal  
  
  cpred <- conformal.glm(fit, parametric = FALSE,  
    nonparametric = TRUE, alpha = alpha,  
    bins = bins, cores = cores)  
  nonparaCI <- cpred$nonparaconformal  
  
  ## obtain LS conformal prediction region  
  pl.tibs <- conformal.pred(x = cbind(x, x^2, x^3), y = y,  
    x0 = cbind(x, x^2, x^3), train.fun = train.fun,  
    predict.fun = predict.fun, alpha = alpha,  
    grid.method = "linear")  
  LSCI <- cbind(pl.tibs$lo, pl.tibs$up)  
  
  ## obtain LSLW conformal prediction region
```

```

cubic.model <- lm(y ~ x + I(x^2) + I(x^3))
abs.resid <- abs(cubic.model$resid)
smooth.call <- smooth.spline(x, abs.resid,
  nknots = 10)
lambda <- smooth.call$lambda
df <- smooth.call$df
mad.train.fun <- function(x, y, out = NULL){
  smooth.spline(x[, 1], y, lambda = lambda,
    df = df, nknots = 10)
}
p2.tibs <- conformal.pred(x = cbind(x, x^2, x^3), y = y,
  x0 = cbind(x, x^2, x^3),
  train.fun = train.fun, predict.fun = predict.fun,
  mad.train.fun = mad.train.fun,
  mad.predict.fun = mad.predict.fun,
  alpha = alpha)
LSLWCI <- cbind(p2.tibs$lo, p2.tibs$up)

## create plots
ix <- sort(x, index.return = TRUE)$ix
par(mfrow = c(1,4), oma = c(4,4,0,0), mar = c(1,1,1,1))

## parametric conformal
plot.new()
plot.window(xlim = c(0, 1), ylim = c(min(y)-1, max(y)+1))
points(x, y, pch = 19, col = rgb(0,0,0,alpha=0.2))
lines(x[ix], paraCI[ix, 1], col = "red", lwd = 1.5)
lines(x[ix], paraCI[ix, 2], col = "red", lwd = 1.5)
axis(1); axis(2)

## LSLW conformal
plot.new()
plot.window(xlim = c(0, 1), ylim = c(min(y)-1, max(y)+1))
points(x, y, pch = 19, col = rgb(0,0,0,alpha=0.2))
lines(x[ix], LSLWCI[ix, 1], col = "red", lwd = 1.5)
lines(x[ix], LSLWCI[ix, 2], col = "red", lwd = 1.5)
axis(1)

## nonparametric conformal
plot.nonparametric(region = nonparaCI, x = x,
  y = y, bins = bins)
axis(1)

## LS conformal
plot.new()
plot.window(xlim = c(0, 1), ylim = c(min(y)-1, max(y)+1))
points(x, y, pch = 19, col = rgb(0,0,0,alpha=0.2))

```

```

lines(x[ix], LSCI[ix, 1], col = "red", lwd = 1.5)
lines(x[ix], LSCI[ix, 2], col = "red", lwd = 1.5)
axis(1)

mtext("x", side = 1, cex = 1, line = 2, outer = TRUE)
mtext("y", side = 2, cex = 1, line = 2, outer = TRUE)
}

```

```

beta <- c(0.5, 1)
makeplots.misspec(n = 150, shape = 5, bins = 2, beta = beta)

```




Figure 29: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 5$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.misspec(n = 150, shape = 10, bins = 2, beta = beta)
```




Figure 30: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 10$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.misspec(n = 150, shape = 20, bins = 2, beta = beta)
```




Figure 31: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 20$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.misspec(n = 150, shape = 30, bins = 2, beta = beta)
```




Figure 32: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 30$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.misspec(n = 150, shape = 40, bins = 2, beta = beta)
```




Figure 33: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 40$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.misspec(n = 150, shape = 50, bins = 2, beta = beta)
```




Figure 34: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 50$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.misspec(n = 150, shape = 60, bins = 2, beta = beta)
```




Figure 35: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 60$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).


```
makeplots.misspec(n = 150, shape = 70, bins = 2, beta = beta)
```




Figure 36: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 70$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.misspec(n = 150, shape = 80, bins = 2, beta = beta)
```




Figure 37: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 80$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.misspec(n = 150, shape = 90, bins = 2, beta = beta)
```




Figure 38: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 90$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.misspec(n = 150, shape = 100, bins = 2, beta = beta)
```




Figure 39: Sim setting: Data is generated as Gamma with $n = 150$, $\text{shape} = 100$, and $\text{bins} = 2$. A cubic regression model is fit to this data for the parametric, LSLW, and LS conformal prediction regions. The regions are depicted as follows: misspecified parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

7.3 Plots corresponding to Section 6

The following function makes the conformal prediction region plots for our linear regression simulations under correct parametric, LS, and LSLW conformal model specification.

```
makeplots.regression <- function(n = 500, bins = 3, sd = 1, beta){  
  
  ## generate random data  
  set.seed(13)  
  alpha <- 0.10  
  cores <- 6  
  p <- k <- length(beta) - 1  
  x <- matrix(runif(n), ncol = p)  
  mu <- (cbind(1, x) %*% beta)  
  y <- rnorm(n = n, mean = mu, sd = sd)  
  y <- y / sd(y)  
  data <- data.frame(y = y, x = x)  
  colnames(data)[2:(p+1)] <- paste("x", 1:p, sep = "")  
  
  ## fit the regression model  
  fit <- glm(y ~ x1, family = "gaussian", data = data)  
  formula <- fit$formula  
  newdata <- data  
  respname <- all.vars(formula)[1]  
  newdata <- newdata[, !(colnames(data) %in% respname)]  
  newdata <- as.matrix(newdata)  
  
  ## obtain parametric and nonparametric conformal  
  ## prediction regions  
  cpred <- conformal.glm(fit, parametric = TRUE,  
    nonparametric = FALSE, alpha = alpha,  
    bins = bins, cores = cores)  
  paraCI <- cpred$paraconformal  
  
  cpred <- conformal.glm(fit, parametric = FALSE,  
    nonparametric = TRUE, alpha = alpha,  
    bins = bins, cores = cores)  
  nonparaCI <- cpred$nonparaconformal  
  
  ## obtain LS conformal prediction region  
  pl.tibs <- conformal.pred(x = x, y = y, x0 = x,  
    train.fun = train.fun,  
    predict.fun = predict.fun, alpha = alpha,  
    grid.method = "linear")  
  LSCI <- cbind(pl.tibs$lo, pl.tibs$up)  
  
  ## obtain LSLW conformal prediction region  
  fit <- lm(y ~ x1, data = data)
```

```

abs.resid <- abs(fit$resid)
smooth.call <- smooth.spline(x, abs.resid,
  nknots = 10)
lambda <- smooth.call$lambda
df <- smooth.call$df
mad.train.fun <- function(x, y, out = NULL){
  smooth.spline(x[, 1], y, lambda = lambda,
    df = df, nknots = 10)
}
p2.tibs <- conformal.pred(x = x, y = y, x0 = x,
  train.fun = train.fun, predict.fun = predict.fun,
  mad.train.fun = mad.train.fun,
  mad.predict.fun = mad.predict.fun, alpha = alpha)
LSLWCI <- cbind(p2.tibs$lo, p2.tibs$up)

## create plots
ix <- sort(x, index.return = TRUE)$ix
par(mfrow = c(1,4), oma = c(4,4,0,0), mar = c(1,1,1,1))

## parametric conformal
plot.new()
plot.window(xlim = c(0, 1), ylim = c(min(y)-1, max(y)+1))
points(x, y, pch = 19, col = rgb(0,0,0,alpha=0.2))
lines(x[ix], paraCI[ix, 1], col = "red", lwd = 1.5)
lines(x[ix], paraCI[ix, 2], col = "red", lwd = 1.5)
axis(1); axis(2)

## LSLW conformal
plot.new()
plot.window(xlim = c(0, 1), ylim = c(min(y)-1, max(y)+1))
points(x, y, pch = 19, col = rgb(0,0,0,alpha=0.2))
lines(x[ix], LSLWCI[ix, 1], col = "red", lwd = 1.5)
lines(x[ix], LSLWCI[ix, 2], col = "red", lwd = 1.5)
axis(1)

## nonparametric conformal
plot.nonparametric(region = nonparaCI, x = x,
  y = y, bins = bins)
axis(1)

## LS conformal
plot.new()
plot.window(xlim = c(0, 1), ylim = c(min(y)-1, max(y)+1))
points(x, y, pch = 19, col = rgb(0,0,0,alpha=0.2))
lines(x[ix], LSCI[ix, 1], col = "red", lwd = 1.5)
lines(x[ix], LSCI[ix, 2], col = "red", lwd = 1.5)
axis(1)

```

```

mtext("x", side = 1, cex = 1, line = 2, outer = TRUE)
mtext("y", side = 2, cex = 1, line = 2, outer = TRUE)
}

```

```

beta <- c(2, 5)
makeplots.regression(n = 150, sd = 1, bins = 2, beta = beta)

```




Figure 40: Sim setting: Data is generated from a normal regression model with constant normal errors with $sd = 1$, $n = 150$, and $bins = 2$. The correct regression model is fit to this data. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.regression(n = 250, sd = 1, bins = 3, beta = beta)
```




Figure 41: Sim setting: Data is generated from a normal regression model with constant normal errors with $sd = 1$, $n = 250$, and $bins = 3$. The correct regression model is fit to this data. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

```
makeplots.regression(n = 500, sd = 1, bins = 3, beta = beta)
```




Figure 42: Sim setting: Data is generated from a normal regression model with constant normal errors with $sd = 1$, $n = 500$, and $bins = 3$. The correct regression model is fit to this data. The regions are depicted as follows: parametric conformal prediction region (far-left panel), LSLW conformal prediction region (middle-left panel), nonparametric conformal prediction region (middle-right panel), LS conformal prediction region (far-right panel).

8 Depiction of conformal prediction regions that appears in Eck et al. [2019]

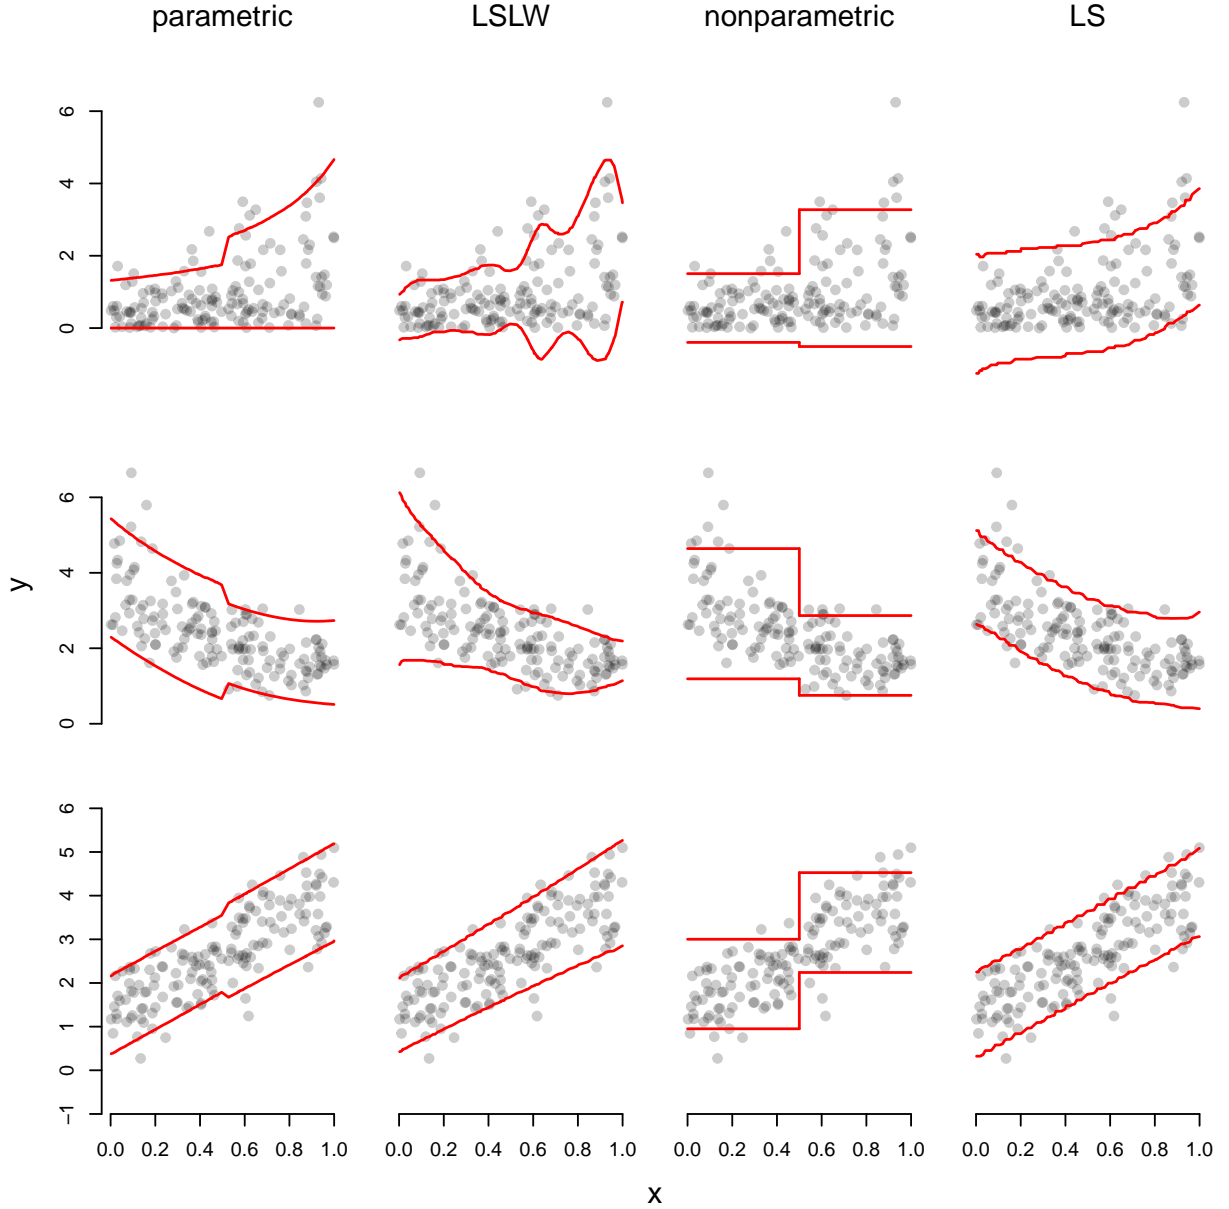


Figure 43: The depiction of conformal prediction regions when $n = 150$ that appears in Eck et al. [2019]. The rows display conformal prediction regions across simulation settings. The columns display the different conformal prediction regions. The top, middle, and bottom rows correspond to simulation setting a with shape parameter equal to 2, simulation setting b with shape parameter equal to 10, and simulation setting c respectively. The first column displays the parametric conformal prediction region which is misspecified in row 2, the second column displays the least squares locally weighted conformal prediction region, the third column displays the nonparametric conformal prediction region, and the fourth column displays the least squares conformal prediction region.

9 Creating this Document

The purpose of this document is to provide a completely reproducible exploration and motivation of conformal prediction. All of the R code presented in this document (or the corresponding .Rnw file) is run when this document is compiled using the linux terminal.

This document is created from its source file `supplement-conformal.Rnw` using `knitr` and the `pdflatex` command in the linux terminal. The `knitr` R package needs to be installed before this document can be compiled. Open R and install this package if you have not previously done so. To compile this document in the linux terminal, enter the command

```
Rscript -e "library(knitr); knitr('supplement-conformal.Rnw')"
```

This produces the LaTeX .tex file with the name `supplement-conformal.tex`. To get a corresponding .pdf file, enter the command

```
pdflatex supplement-conformal.tex
```

You may want to run the previous command twice in order to get labels and references right.

References

Daniel J. Eck. *conformal.glm: Conformal Prediction for Generalized Linear Regression Models*, 2018. <https://github.com/DEck13/conformal.glm>.

Daniel J. Eck, Forrest W. Crawford, and Peter M. Aronow. Conformal inference for exponential families and generalized linear models. *preprint*, 2019.

Jing Lei and Larry Wasserman. Distribution-free prediction bands for nonparametric regression. *Journal of the Royal Statistical Society series B*, 76, 1:71–96, 2014.

Jing Lei, Max G. Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523):1094–1111, 2018.

Mike Meredith and John Kruschke. *HDInterval: Highest (Posterior) Density Intervals*, 2018. <https://cran.r-project.org/web/packages/HDInterval/>.

Ryan Tibshirani. *conformalInference: Tools for conformal inference in regression*, 2016. <https://github.com/ryantibs/conformal>.